

Reproducible Polyglot Data Science

Bruno Rodrigues

2025-12-31

Table of contents

Welcome!	1
A modern, unified, and language-agnostic workflow for data science using Nix.	1
Preface	5
1 Introduction	9
1.1 Who is this book for?	9
1.2 What is the aim of this book?	10
1.3 Prerequisites	12
1.4 What actually is reproducibility?	14
1.4.1 Using open-source tools is a hard require- ment	14
1.4.2 Hidden dependencies can hinder repro- ducibility	17
1.4.3 The requirements of a RAP	18
1.4.4 Are there different types of reproducibility?	19
2 The Nix Package Manager	25
2.1 Introduction	25
2.2 Important Concepts	27
2.2.1 Derivations	28
2.2.2 Dependencies of derivations	29
2.2.3 The Nix store and hermetic builds	30
2.2.4 Other key Nix concepts	31
2.3 Caveats	32

Table of contents

2.4	In summary	34
3	—	37
4	Reproducibility with Nix	39
4.1	Why Reproducibility? Why Nix? (<i>1h30</i>)	39
4.1.1	Motivation: Reproducibility in Scientific and Data Workflows	39
4.1.2	Problems with Ad-Hoc Tools	40
4.1.3	Nix, a declarative package manager	40
4.1.4	The <code>rix</code> package	41
4.1.5	Installing Nix	42
4.1.6	Temporary shells	45
4.2	Session 1.2 – Dev Environments with Nix (<i>1h30</i>)	46
4.2.1	Some Nix concepts	46
4.2.2	Derivations	47
4.2.3	Using <code>{rix}</code> to generate development environments	48
4.2.4	Using <code>nix-shell</code> to Launch Environments	50
4.2.5	Pinning with <code>nixpkgs</code>	50
4.2.6	Installing Python packages not available via <code>nixpkgs</code> (impure)	51
4.3	Configuring your IDE	54
4.3.1	Pre-requisites	54
4.3.2	<code>direnv</code>	56
4.3.3	In summary and next steps	58
4.4	Starting a new project:	59
4.5	Hands-On Exercises	60
5	Reproducible Development Environments with rix	61
5.1	Introduction	61
5.1.1	The Reproducibility Challenge	61
5.1.2	Component Closures: The Nix Approach	62
5.1.3	The Polyglot Challenge	62

Table of contents

5.1.4	Enter rix	63
5.2	The rix() Function	63
5.3	Choosing an R Version or Date	65
5.3.1	Using dates vs versions	66
5.3.2	The rstats-on-nix fork	66
5.4	Installing R Packages	67
5.4.1	From CRAN/Bioconductor	67
5.4.2	Installing archived versions	67
5.4.3	Installing from GitHub	68
5.4.4	Installing local packages	69
5.4.5	Why NOT to use install.packages()	69
5.5	System Tools and TexLive	70
5.5.1	Adding system packages	70
5.5.2	TexLive for literate programming	70
5.5.3	Python and Julia integration	71
5.5.4	Installing Python packages with uv (im-pure)	72
5.6	Building and Using Environments	74
5.6.1	Building with nix-build	74
5.6.2	Entering with nix-shell	74
5.6.3	Running scripts	75
5.6.4	Pure shells for complete isolation	75
5.6.5	Garbage collection	76
5.7	Converting renv Projects	76
5.7.1	Recommended workflow	76
5.7.2	Caveats	77
5.8	Summary	77
6	Functional Programming	79
6.1	Introduction: From Scripts to Functions	79
6.1.1	The Notebook Problem	80
6.1.2	The Functional Solution	81
6.1.3	FP vs OOP: Transformations vs Actors	81
6.1.4	Why Does This Matter for Data Science?	82

Table of contents

6.2	Purity and Side Effects	83
6.2.1	Handling “Impure” Operations like Randomness	84
6.2.2	The OOP Caveat in Python	85
6.3	Writing Your Own Functions	86
6.3.1	In R	86
6.3.2	In Python	88
6.4	The Functional Toolkit: Map, Filter, and Reduce	88
6.4.1	1. Mapping: Applying a Function to Each Element	88
6.4.2	2. Filtering: Keeping Elements That Match a Condition	90
6.4.3	3. Reducing: Combining All Elements into a Single Value	91
6.5	The Power of Composition	92
6.5.1	The Composition Challenge in Python . .	93
6.6	Summary	93
7	Building Reproducible Pipelines with rixpress	95
7.1	Introduction: From Scripts and Notebooks to Pipelines	95
7.1.1	The Evolution of Build Automation	95
7.1.2	The Separation Problem	96
7.1.3	The Imperative Approach: Make + Docker	96
7.1.4	rixpress: Unified Orchestration	99
7.2	What is rixpress?	101
7.3	Getting Started	103
7.3.1	Initialising a project	103
7.3.2	Defining the environment	103
7.3.3	Defining the pipeline	104
7.4	Core Functions	105
7.4.1	Defining derivations	105
7.4.2	Building the pipeline	106
7.4.3	Inspecting outputs	106

7.4.4	Visualizing the pipeline	106
7.5	Polyglot Pipelines	107
7.5.1	Method 1: Using language converters . .	108
7.5.2	Method 2: Using universal data formats .	108
7.6	Caching and Incremental Builds	109
7.7	Build Logs and Debugging	110
7.8	ryxpress: The Python Interface	111
7.9	Running Someone Else’s Pipeline	112
7.10	Exporting and Importing Artifacts	112
7.11	Real-World Examples	113
7.11.1	Example 1: Machine Learning with XG-Boost	113
7.11.2	Example 2: Reading Many Input Files . .	116
7.11.3	Example 3: Full Python→R→Quarto Workflow	117
7.11.4	More Examples	119
7.12	Summary	120
8	Containerisation With Docker	121
8.1	Introduction	121
8.1.1	Spatial vs Temporal Reproducibility . .	121
8.1.2	The Best of Both Worlds	122
8.1.3	Interactive Development vs Distribution .	123
8.2	Docker Essentials	123
8.2.1	Installing Docker	123
8.2.2	The Rocker Project and Image Registries .	125
8.2.3	Basic Docker Workflow	127
8.3	Making Our Own Images	129
8.3.1	A Minimal Dockerfile with Nix	129
8.3.2	Splitting Into a Reusable Base Image . .	131
8.4	Publishing Images on Docker Hub	133
8.4.1	Sharing Without Docker Hub	134
8.5	What If You Don’t Use Nix?	134

Table of contents

8.6	Building Docker Images Directly with Nix	135
8.6.1	Why Skip the Dockerfile?	136
8.6.2	A Basic Example	136
8.6.3	Benefits Over Dockerfiles	138
8.6.4	Layered Images for Efficiency	138
8.6.5	When to Use This Approach	139
8.7	Dockerizing a rixpress Pipeline	139
8.7.1	Step 1: The Dockerfile	140
8.7.2	Step 2: The Export Script	140
8.7.3	Step 3: Build and Run	141
8.8	Summary	142
8.9	Further Reading	142
9	Robust Pipelines with Monads	143
9.1	The Problem: Decorated Functions Don't Compose	143
9.2	The Solution: Function Factories and Bind	145
9.3	The chronicler Package	146
9.4	The Maybe Monad: Handling Errors	147
9.5	Monads in Python: talvez	148
9.6	Building Robust Pipelines	149
9.7	Integrating with rixpress	151
9.8	Python Integration with talvez	153
9.9	Summary	154
References		155

Welcome!

A modern, unified, and language-agnostic workflow for data science using Nix.

This book is a complete reimagining of my previous work, “Building Reproducible Analytical Pipelines with R.” If you’re looking for that book, you can find it here. But if you’re ready for the next step, you’re in the right place.

Data scientists, statisticians, analysts, researchers, and many other professionals write a lot of code.

Not only do they write a lot of code, but they must also read and review a lot of code as well. They either work in teams and need to review each other’s code, or need to be able to reproduce results from past projects, be it for peer review or auditing purposes. And yet, they never, or very rarely, get taught the tools and techniques that would make the process of writing, collaborating, reviewing and reproducing projects possible.

Which is truly unfortunate because software engineers face the same challenges and solved them decades ago.

The aim of this book is to teach you how to use some of the best practices from software engineering and DevOps to make your

Welcome!

projects robust, reliable and reproducible. It doesn't matter if you work alone, in a small or in a big team. It doesn't matter if your work gets (peer-)reviewed or audited: the techniques presented in this book will make your projects more reliable and save you a lot of frustration!

As someone whose primary job is analysing data, you might think that you are not a developer. It seems as if developers are these genius types that write extremely high-quality code and create these super useful packages. The truth is that you are a developer as well. It's just that your focus is on writing code for your purposes to get your analyses going instead of writing code for others. Or at least, that's what you think. Because in others, your team-mates are included. Reviewers and auditors are included. Any people that will read your code are included, and there will be people that will read your code. At the very least future you will read your code. By learning how to set up projects and write code in a way that future you will understand and not want to murder you, you will actually work towards improving the quality of your work, naturally.

The book can be read for free on <https://b-rodrigues.github.io/reproducible-data-science/> and you'll be able buy a DRM-free Epub or PDF on Leanpub¹ once there's more content.

This book is the culmination of my previous works. I started by writing a book focused on R, and then began working on a Python edition. During that process, I had a realization: tackling reproducibility one language at a time was solving the symptoms, not the root cause. The real solution needed to be universal, powerful, and capable of handling any language or tool we might need.

That universal solution is Nix.

¹<https://leanpub.com/>

This book moves beyond language-specific tooling. It presents a holistic workflow where R, Python, and Julia are not competitors, but collaborators in a single, cohesive, and perfectly reproducible environment. We will cover:

- The Nix Philosophy: Why Nix is the ultimate tool for solving the “it works on my machine” problem, once and for all.
- Declarative Environments with `{trix}`: How to use a simple R interface to define exact, bit-for-bit reproducible software environments that include specific versions of R, Python, Julia, their packages, and any system-level dependencies.
- Polyglot Pipelines with `{trixpress}` (R) or `ryxpress` (Python): How to orchestrate complex analytical pipelines that seamlessly pass data between different languages, all managed by the Nix build system.
- Unit Testing and Functional Programming: Core principles for writing robust, testable, and maintainable code, no matter the language.
- Distribution and Automation: How to package your entire reproducible pipeline into a Docker container for easy sharing and automate your workflow with GitHub Actions.

While this is not a book for beginners (you should be familiar with at least one data-centric programming language before reading this), I will not assume that you have any knowledge of the tools discussed. But be warned, this book will require you to take the time to read it, and then type on your computer. Type *a lot*.

I hope that you will enjoy reading this book and applying the ideas in your day-to-day, ideas which hopefully should improve the reliability, traceability and reproducibility of your code.

Welcome!

If you find this book useful, don't hesitate to let me know! You can submit issues, suggest improvements, and ask questions on the book's Github repository.

If you want to get to know me better, read my bio².

²<https://www.brodrigues.co/about/me/>

Preface

Three years ago, I wrote a book with a straightforward premise: by borrowing a few key ideas from software engineering, people who analyse data could save themselves a great deal of frustration. The response to that book was more positive than I could have ever hoped for, and it confirmed a suspicion I had: we, as a community, are hungry for better ways to work.

That book, however, focused exclusively on the R ecosystem. A recurring question I received was, “This is great, but what about Python?” It was a fair question. The world of data science is not a monologue; it’s a conversation between languages. So, I began what felt like the logical next step: writing a Python edition.

I mapped out the chapters and identified the equivalent tools, `pipenv` for dependency management, `ploomber` for pipelines, and started writing. But as I went deeper, a nagging feeling grew. I was solving the same problems all over again, just with a different set of tools. This feeling was compounded by the rapid churn within the Python ecosystem itself. How many package managers have been created to solve virtual environment management? As of writing, `uv` is all the rage, and while it may be here to stay, history suggests a new contender is always just around the corner.

This pointed to a larger issue. I am convinced that the future of data science is polyglot. An R user and a Python user, both following my original advice, would end up with reproducible

Preface

projects, but their workflows would be fundamentally incompatible. They couldn't easily share an environment or build a single pipeline that leverages the strengths of both languages. While companies like Posit have made excellent progress in making it easier to call Python from R, setting up a truly integrated development environment remains a challenge. And what if you wish to bring Julia, the other language of data analysis, into the fold? It is not as popular as Python or R, but it has its own distinct appeal and advantages.

I realised I was treating the symptoms, not the disease. The root problem wasn't "How do I make R reproducible?" or "How do I make Python reproducible?". The real challenge was the lack of a universal foundation that could handle *any* language, *any* tool, and *any* system dependency with absolute, bit-for-bit precision.

That's when I stopped writing the Python book.

The solution wasn't to create another language-specific guide but to find a tool that operated at a more fundamental level. That tool, I am now convinced, is **Nix**.

Nix is not just another package manager; it is a powerful, declarative system for building and managing software environments. It allows us to define the *entire* computational environment—from the operating system libraries up to the specific versions of our R and Python packages—in a single, simple text file. When you use Nix, the phrase "it works on my machine" becomes obsolete. It is replaced by the guarantee: "it builds identically, everywhere, every time"—with a few caveats that we will explore, of course.

This book is the result of that realisation. It is a complete reimaging of the original. We are moving away from language-specific patchworks and toward a unified, polyglot workflow. We

will use Nix as our bedrock, with the `{rix}` and `{rixpress}` R packages (or `ryxpress` for Python) serving as our friendly interface to its power. You will learn to build pipelines where R, Python, and Julia aren't just neighbours; they are collaborators, working together in a single, perfectly reproducible environment.

A note for Python-first users: do not be deterred by the fact that `{rix}` and `{rixpress}` are R packages: there is a Python version called `ryxpress` that will allow you to run your pipelines from an interactive Python session. You will be able to use them to define your environments (even integrating tools like `uv`) and orchestrate your pipelines, while doing all of your analytical work exclusively in Python. In this workflow, R simply becomes a convenient configuration language.

The core message from three years ago remains unchanged. You, as someone who writes code to analyse data, are a developer. Your work is important, and it deserves to be reliable. This book aims to give you the tools and the mindset to achieve that. The journey is more ambitious this time, but the payoff is far greater.

I hope you'll join me.

You can read this book for free online at <https://b-rodrigues.github.io/reproducible-data-science/>.

You can submit issues, suggest improvements, and ask questions on the book's Github repository.

1 Introduction

Just like the previous, R-focused edition of this book, this one will not teach you about machine learning, statistics, or visualisation.

The goal is to teach you a set of tools, practices, and project management techniques that should make your projects easier to reproduce, replicate, and retrace, with a focus on polyglot (multilingual) projects. I believe that with LLMs, polyglot projects will become increasingly common.

Before LLMs, if you were a Python user, you would avoid R like the plague, and vice versa. This is understandable: even though both are high-level scripting languages, and an R user could likely read and understand Python code (and vice versa), it is still a pain in the loins to set up another language just to run a few lines of code. Now, with LLMs, you can have a model generate the code, and depending on what you're doing, it's likely to be correct. However, setting up another language is still quite annoying. This is where Nix comes in. Nix makes adding another language to a project extremely simple.

1.1 Who is this book for?

This book is for anyone who uses raw data to build any type of output. This could be a simple quarterly report, in which

1 Introduction

data is used for tables and graphs, a scientific article for a peer-reviewed journal, or even an interactive web application. The specific output doesn't matter, because the process is, at its core, always very similar:

- Get the data;
- Clean the data;
- Write code to analyse the data;
- Put the results into the final product.

This book assumes some familiarity with programming, particularly with the R and Python languages. I will not discuss Julia in great detail, as I am not familiar enough with it to do it justice. That being said, I will show you how to add Julia to a project and use it effectively if you need to.

1.2 What is the aim of this book?

The aim of this book is to make the process of analysing data as reliable, retraceable, and reproducible as possible, and to do this by design. This means that once you're done with the analysis, you're done. You don't want to spend time, which you often don't have anyway, to rewrite or refactor an analysis to make it reproducible after the fact. We both know this is not going to happen. Once an analysis is finished, it's on to the next one. If you need to rerun an older analysis (for example, because the data has been updated), you'll simply figure it out at that point, right? That's a problem for Future You. Hopefully, Future You will remember every quirk of your code, know which script to run at which point, which comments are outdated, and what features of the data need to be checked... You had better hope Future You is a more diligent worker than you are!

1.2 What is the aim of this book?

Going forward, I'm going to refer to a project that is reproducible as a “Reproducible Analytical Pipeline”, or RAP for short. There are only two ways to create a RAP: either you are lucky enough to have someone on your team whose job is to turn your messy code into a RAP, or you do it yourself. The second option is by far the most common. The issue, as stated above, is that most of us simply don't do it. We are always in a rush to get to the results and don't think about making the process reproducible, because we assume it takes extra time that is better spent on the analysis itself. This is a misconception, for two reasons.

The first is that employing the techniques we will discuss in this book won't actually take much time. As you will see, they are not things you “add on top of” the analysis, but are part of the analysis itself, and they will also help with managing the project. Some of these techniques, especially testing, will even save you time and headaches.

The second reason is that an analysis is never a one-shot. Only the simplest tasks, like pulling a single number from a database, might be. Even then, chances are that once you provide that number, you'll be asked for a variation of it (for example, disaggregated by one or several variables). Or perhaps you'll be asked for an update in six months. You will quickly learn to keep that SQL query in a script somewhere to ensure consistency. But what about more complex analyses? Is keeping the script enough? It is a good start, of course, but very often, there is no single script, or a script for each step of the analysis is missing.

I've seen this play out many times in many different organisations. It's that time of the year again, and a report needs to be written. Ten people are involved, and just gathering the data

1 Introduction

is already complicated. Some get their data from Word documents attached to emails, some from a website, some from a PDF report from another department. I remember a story a senior manager at my previous job used to tell: once, a client put out a call for a project that involved setting up a PDF scraper. They periodically needed data from another department that only came in PDFs. The manager asked what was, at least from our perspective, an obvious question: “Why can’t they send you the underlying data in a machine-readable format?” They had never thought to ask. So, my manager went to that department and talked to the people putting the PDF together. Their answer? “Well, we could send them the data in any format they want, but they’ve asked for the tables in a PDF.”

So the first, and probably most important, lesson here is: when starting to build a RAP, make sure you talk with all the people involved.

1.3 Prerequisites

You should be comfortable with the command line. This book will not assume any particular Integrated Development Environment (IDE), so most of what I’ll show you will be done via the command line. That said, I will spend some time helping you set up a data science-focused IDE, Positron, to work seamlessly with this workflow. The command line may be over 50 years old, but it is not going anywhere. In fact, thanks to the rise of LLMs, it seems to be enjoying a resurgence. Since these models generate text, it is far simpler to ask one for a shell command to solve a problem than to have it produce detailed instructions on where to click in a graphical user interface (GUI). Knowing your way around the command line is also essential for working with

modern data science infrastructure: continuous integration platforms, Docker, remote servers... they all live in the terminal. So, if you're not at all familiar with the command line, you might need to brush up on the basics. Don't worry, though; this isn't a book about the intricacies of the Unix command line. The commands I'll show you will be straightforward and directly applicable to the task at hand.

This means however that if you are using Windows, first of all, why? and second of all, you will have to set up Windows Subsystem for Linux. This is because there is no native Nix implementation for Windows, and so we need to run the Linux version through WSL. Don't worry though, it's not that hard, and you can then use an IDE from Windows to work with the environments managed by Nix, in a very seamless way (but seriously, consider, if you can, switching to Linux. How can you tolerate ads (ADS!) in the start menu).

Ideally, you should be comfortable with either R or Python. This book will assume that you have been using at least one of these languages for some projects and want to improve how you manage complex projects. You should know about packages and how to install them, have written some functions, understand loops, and have a basic knowledge of data structures like lists. While this is not a book on visualisation, we will be making some graphs as well.

Our aim is to write code that can be executed non-interactively. This is because a necessary condition for a workflow to be reproducible and to qualify as a RAP is for it to be executed by a machine, automatically, without any human intervention. This is the second lesson of building RAPs: there should be no human intervention needed to get the outputs once the RAP has started. If you achieve this, then your workflow is likely reproducible, or can at least be made so much more easily than if

1 Introduction

it requires special manipulation by a human somewhere in the loop.

1.4 What actually is reproducibility?

A reproducible project means that it can be rerun by anyone at zero (or very minimal) cost. But there are different levels of reproducibility, which I will discuss in the next section. Let's first outline the requirements that a project must meet to be considered a RAP.

1.4.1 Using open-source tools is a hard requirement

Open source is a hard requirement for reproducibility. No ifs, ands, or buts. I'm not just talking about the code you wrote for your research paper or report; I'm talking about the entire ecosystem you used to write your code and build the workflow.

Is your code open? Good. Or is it at least available to others in your organisation, in a way that they could re-execute it if needed? Also good.

But is it code written in a proprietary program, like STATA, SAS, or MATLAB? Then your project is not reproducible. It doesn't matter if the code is well-documented, well-written, and available on a version control system. The project is simply not reproducible. Why?

Because on a long enough time horizon, there is no way to re-execute your code with the exact same version of the proprietary language and operating system that were used when the project

1.4 What actually is reproducibility?

was developed. As I'm writing these lines, MATLAB, for example, is at version R2025a. Buying an older version may not be simple. I'm sure if you contact their sales department, they might be able to sell you an older version. Maybe you can even re-download older versions you've already purchased from their website. But maybe it's not that straightforward. Or maybe they won't offer this option in the future. In any case, if you search for "purchase old version of Matlab," you will see that many researchers and engineers have this need. :::{.content-hidden when-format="pdf"}

Wanting to run older versions of analytics software is a recurrent need.

:::

1 Introduction

Old version of matlab

⊕ Follow

4 views (last 30 days)

 [REDACTED] on 29 Nov 2018
STAFF MVP on 29 Nov 2018

Vote | 1 Link

Hallo after a few years we need to use again an old program written with matlab R12 6.0.0.88. We don't find the installation CD, can we buy again this old version of the program? Thanks best regard

0 Comments

[Sign in to comment.](#)

[Sign in to answer this question.](#)

Answers (1)

 [REDACTED] STAFF MVP on 29 Nov 2018

Vote | 0 Link

Have you tried running the old program on a more recent release of MATLAB?

MATLAB 6.0 (R12) is [eighteen years old](#) (released in November 2000) and I think it highly unlikely you'll be able to get it working on a new operating system. The [Windows system requirements](#) lists several Windows versions on which that release was supported, the [newest](#) of which was Windows ME which was released in September 2000. Microsoft ended mainstream support for this OS in 2003 and ended extended support in July 2006 according to [Wikipedia](#).

0 Comments

[Sign in to comment.](#)

Figure 1.1: Wanting to run older versions of analytics software is a recurrent need.

And if you're running old code written for version, say, R2008a, there's no guarantee that it will produce the exact same results on version R2025a. That's without even mentioning the toolboxes (if you're not familiar with them, they're MATLAB's equivalent of packages or libraries). These evolve as well, and there's no guarantee that you can purchase older versions. It's also likely that newer toolboxes cannot even run on older versions of MATLAB.

Let me be clear: what I'm describing here for MATLAB could also be said for any other proprietary programs still commonly (and unfortunately) used in research and statistics, like STATA,

SAS, or SPSS. Even if some of these vendors provide ways to run older versions of their software, the fact that you have to rely on them for this is a barrier to reproducibility. There is no guarantee they will provide this option forever. Who can guarantee that these companies will even be around forever? More likely, they might shift from a program you install on your machine to a subscription-based model.

For just 199€ a month, you can execute your SAS (or whatever) scripts on the cloud! Worried about data confidentiality? No problem, data is encrypted and stored safely on our secure servers! Run your analysis from anywhere and don't worry about your cat knocking coffee over your laptop! And if you purchase the pro licence, for an additional 100€ a month, you can even execute your code in parallel!

Think this is science fiction? There is a growing and concerning trend for vendors to move to a Software-as-a-Service model with monthly subscriptions. It happened to Adobe's design software, and the primary reason it hasn't yet happened for data analytics tools is data privacy concerns. Once those are deemed "solved," I would not be surprised to see a similar shift.

1.4.2 Hidden dependencies can hinder reproducibility

Then there's another problem. Let's suppose you've written a thoroughly tested and documented workflow and made it available on GitHub (and let's even assume the data is freely available and the paper is open access). Or, if you're in the private sector, you've done all of the above, but the workflow is only available internally.

1 Introduction

Let's further assume that you've used R, Python, or another open-source language. Can this analysis be said to be reproducible? Well, if it ran on a proprietary operating system, then the conclusion is: your project is not fully reproducible.

This is because the operating system the code runs on can also influence the outputs. There are particularities in operating systems that may cause certain things to work differently. Admittedly, this is rarely a problem in practice, but it does happen¹, especially if you're working with high-precision floating-point arithmetic, as you might in the financial sector, for instance.

Thankfully, as you will see in this book, there is no need to change your operating system to deal with this issue.

1.4.3 The requirements of a RAP

So where does that leave us? For a project to be truly reproducible, it has to respect the following points:

- Source code must be available and thoroughly tested and documented (which is why we will be using Git and GitHub).
- All dependencies must be easy to find and install (we will deal with this using dependency management tools).
- It must be written in an open-source programming language (no-code tools like Excel are non-reproducible by default because they can't be used non-interactively, which is why we will be using languages like R, Python, and Julia).

¹<https://github.com/numpy/numpy/issues/9187>

- The project needs to run on an open-source operating system (we can deal with this without having to install and learn a new OS, thanks to tools like Docker).
- The data and the final report must be accessible—if not publicly, then at least within your company. This means the concept of “scripts and/or data available upon request” belongs in the bin.

Availability of data

Data available upon reasonable request

Figure 1.2: A real sentence from a real paper published in *THE LANCET Regional Health*. How about *make the data available and I won't scratch your car*, how's that for a reasonable request?

1.4.4 Are there different types of reproducibility?

Let's take one step back. We live in the real world, where constraints outside of our control can make it impossible to build a true RAP. Sometimes we need to settle for something that might not be perfect, but is the next best thing.

In what follows, let's assume the code is tested and documented, so we will only discuss the pipeline's execution.

The *least* reproducible pipeline would be something that works, but only on your machine. This could be due to hardcoded paths

1 Introduction

that only exist on your laptop. Anyone wanting to rerun the pipeline would need to change them. This should be documented in a README, which we've assumed is the case. But perhaps the pipeline only runs on your laptop because the computational environment is hard to reproduce. Maybe you use software, even open-source software, that is not easy to install (anyone who has tried to install R packages on Linux that depend on `{rJava}` knows what I'm talking about).

A better, though still imperfect, pipeline would be one that could be run more easily on any similar machine. This could be achieved by avoiding hardcoded absolute paths and by providing instructions to set up the environment. For example, in Python, this could be as simple as providing a `requirements.txt` file that lists the project's dependencies, which could be installed using `pip`:

```
pip install -r requirements.txt
```

Doing this helps others (or Future You) install the required packages. However, this is not enough, as other software on your system, outside of what `pip` manages, can still impact the results.

1.4 What actually is reproducibility?

You should also ensure that people run the same analysis on the same versions of R or Python that were used to create it. Just installing the right packages is not enough. The same code can produce different results on different versions of a language, or not run at all. If you've been using Python for some time, you certainly remember the switch from Python 2 to Python 3. Who knows, the switch to Python 4 might be just as painful!

The take-away message is that relying on the language itself being stable over time is not a sufficient condition for reproducibility. We have to set up our code in a way that is *explicitly* reproducible by dealing with the versions of the language itself.

So what does this all mean? It means that reproducibility exists on a continuum. Depending on the constraints you face, your project can be “not very reproducible” or “totally reproducible”. Let’s consider the following list of factors that can influence how reproducible your project truly is:

- Version of the programming language used.
- Versions of the packages/libraries of said programming language.
- The operating system and its version.
- Versions of the underlying system libraries (which often go hand-in-hand with the OS version, but not always).
- And even the hardware architecture that you run the software stack on.

By “reproducibility is on a continuum,” I mean that you can set up your project to take none, one, two, three, four, or all of the preceding items into consideration.

This is not a novel, or new idea. Peng (2011) already discussed this concept but named it the *reproducibility spectrum*:

1 Introduction

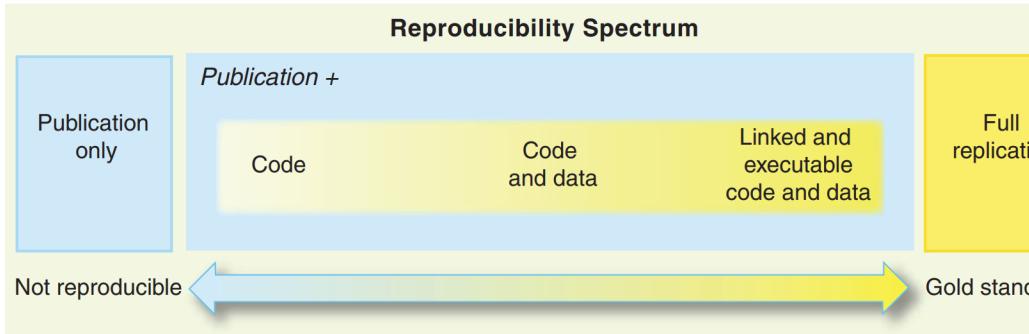


Figure 1.3: The reproducibility spectrum from Peng’s 2011 paper.

Let me finish this introduction by discussing the last item on the list: hardware architecture. In 2020, Apple changed the hardware architecture of their computers. Their new machines no longer use Intel CPUs, but instead Apple’s own proprietary architecture (Apple Silicon) based on the ARM specification. Concretely, this means that binary packages built for Intel-based Apple computers cannot run on their new machines, at least not without a compatibility layer. If you have a recent Apple Silicon Mac and need to install old packages to rerun a project (and we will learn how to do this), they need to be compiled to work on Apple Silicon first. While a compatibility layer called Rosetta 2 exists, my point is that you never know what might come in the future. The ability to compile from source is important because it requires the fewest dependencies outside of your control. Relying on pre-compiled binaries is not future-proof, which is another reason why open-source tools are a hard requirement for reproducibility.

For you Windows users, don’t think that the preceding paragraph does not concern you. It is very likely that Microsoft will

1.4 What actually is reproducibility?

push for OEM manufacturers to build more ARM-based computers in the future. There is already an ARM version of Windows, and I believe Microsoft will continue to support it. This is because ARM is much more energy-efficient than other architectures, and any manufacturer can build its own ARM CPUs by purchasing a license—a very interesting proposition from a business perspective.

It is also possible that we will move towards more cloud-based computing, though I think this is less likely than the hardware shift. In that case, it is quite likely that the actual code will be running on Linux servers that are ARM-based, due to energy and licensing costs. Here again, if you want to run your historical code, you'll have to compile old packages and programming language versions from source.

Ok, this might all seem incredibly complicated. How on earth are we supposed to manage all these risks and balance the immediate need for results with the future need to rerun an old project? And what if rerunning it is never needed?

As you shall see, this is not as difficult as it sounds, thanks to Nix and the tireless efforts of the `nixpkgs` maintainers who work to build truly reproducible packages.

Let's dive in!

2 The Nix Package Manager

2.1 Introduction

Nix is a package manager that can be installed on your computer, regardless of the operating system. If you are familiar with the Ubuntu Linux distribution, you have likely used `apt-get` to install software. On macOS, you may have used `homebrew` for similar purposes. Nix functions in a comparable way but has many advantages over classic package managers, as it focuses on reproducible builds and downloads packages from `nixpkgs`, currently the largest software repository¹.

What makes Nix particularly useful for reproducible projects? Why not use another package manager? Wouldn't we achieve the same thing?

To answer that, let's start from the beginning. To ensure a project is reproducible, you need to deal with at least four challenges:

- Ensure the required version of your programming language (R, Python, etc.) is installed.
- Ensure the required versions of all packages are installed.

¹<https://repology.org/repositories/graphs>

2 The Nix Package Manager

- Ensure all necessary system dependencies are installed (for example, a working Java installation for the `rJava` R package on Linux).
- Ensure you can install all of this on the hardware you have on hand.

The current consensus for tackling the first three points is often a mixture of tools: Docker for system dependencies, `renv` or `uv` for package management, and tools like the R installation manager (`rig`) for language versions. As for the last point, hardware architecture, the only way out is to be able to compile the software for the target platform. This involves a lot of moving parts and requires significant knowledge to get right.

With Nix, we can handle all of these challenges with a single tool.

The first advantage of Nix is that its repository, `nixpkgs`, is humongous. As of this writing, it contains over 120,000 pieces of software, including the *entirety of CRAN and Bioconductor*. This means you can use Nix to handle everything: R, Python, Julia, their respective packages, and any other software available through `nixpkgs`, making it particularly useful for polyglot pipelines.

The second, and most crucial, advantage is that Nix allows you to install software in (relatively) isolated environments. When you start a new project, you can use Nix to install a project-specific version of R and all its packages. These dependencies are used only for that project. If you switch to another project, you switch to a different, independent environment. But this also means that all the dependencies of R and R packages, plus all of their dependencies and so on get installed as well. Your project's development environment will not depend on anything outside of it.

This is similar to `{renv}`, but the difference is profound: you get not only a project-specific library of R packages but also a project-specific R version and all the necessary system dependencies. For example, if you need `{xlsx}`, Nix automatically figures out that Java is required and installs and configures it for you, without any intervention.

What's more, you can *pin* your project to a specific revision of the `nixpkgs` repository. This ensures that every package Nix installs will always be at the exact same version, regardless of when or where the project is built. The environment is defined in a simple plain-text file, and anyone using that file will get a byte-for-byte identical environment, even on a different operating system.

2.2 Important Concepts

Before we start using Nix, it is important to spend some time learning about some Nix-centric concepts, starting with the *derivation*.

In Nix terminology, a derivation is *a specification for running an executable on precisely defined input files to repeatably produce output files at uniquely determined file system paths*.

In simpler terms, a derivation is a recipe with precisely defined inputs, steps, and a fixed output. This means that given identical inputs and build steps, the exact same output will always be produced. To achieve this level of reproducibility, several important measures must be taken:

- All inputs to a derivation must be explicitly declared.

2 The Nix Package Manager

- Inputs include not just data files but also software dependencies, configuration flags, and environment variables—essentially, anything necessary for the build process.
- The build process takes place in a *hermetic* sandbox to ensure the exact same output is always produced.

The next sections explain these three points in more detail.

2.2.1 Derivations

Here is an example of a simple Nix expression:

```
let
  pkgs = import (fetchTarball
    "https://github.com/rstats-on-nix/nixpkgs/archive/2025-04-14.tar.gz"
    {});
in
pkgs.stdenv.mkDerivation {
  name = "filtered_mtcars";
  buildInputs = [ pkgs.gawk ];
  dontUnpack = true;
  src = ./mtcars.csv;
  installPhase = ''
    mkdir -p $out
    awk -F',' 'NR==1 || $9=="1" { print }' $src >
    $out/filtered.csv
  '';
}
```

Without going into too much detail, this code uses `awk`, a common Unix data processing tool, to filter the `mtcars.csv` file. As you can see, a significant amount of boilerplate is required

for this simple operation. However, this approach is completely reproducible: the dependencies are declared and pinned to a specific version of the `nixpkgs` repository. The only thing that could make this small pipeline fail is if the `mtcars.csv` file is not provided to it.

Nix builds the `filtered.csv` in two steps: it first generates a *derivation* from this expression, and only then does it build the output. For clarity, I will refer to code like the example above as a *derivation* rather than an expression, to avoid confusion with the concept of an *expression* in R.

The goal of the tools we will use in this book, `{frx}` and `{frxpress}` (or `ryxpress` if you prefer using Python), is to help you create pipelines from such derivations without needing to learn the Nix language itself, while still benefiting from its powerful reproducibility features.

2.2.2 Dependencies of derivations

Nix requires that the dependencies of any derivation be explicitly listed and managed by Nix itself. If you are building an output that requires Quarto, then Quarto must be explicitly listed as an input, even if you already have it installed on your system. The same applies to Quarto’s dependencies, and their dependencies, all the way down. To run a linear regression with R, you essentially need Nix to build the entire universe of software that R depends on first.

In Nix terms, this complete set of packages is what its author, Eelco Dolstra, refers to as a *component closure*:

The idea is to always deploy component closures: if we deploy a component, then we must also deploy its

2 The Nix Package Manager

dependencies, their dependencies, and so on. That is, we must always deploy a set of components that is closed under the ‘depends on’ relation.

(*Nix: A Safe and Policy-Free System for Software Deployment*, Dolstra et al., 2004).

Figure 4 of Dolstra et al. (2004)

In the figure, `subversion` depends on `openssl`, which itself depends on `glibc`. Similarly, if you write a derivation to filter `mtcars`, it requires an input file, `R`, `{dplyr}`, and all of their respective dependencies. All of these must be managed by Nix. If any dependency exists “outside” this closure, the pipeline will only *work on your machine*—defeating the purpose of reproducibility.

2.2.3 The Nix store and hermetic builds

When building derivations, their outputs are saved into the **Nix store**. Typically located at `/nix/store/`, this folder contains all the software and build artefacts produced by Nix.

For example, the output of a derivation might be stored at a path like `/nix/store/81k4s9q652j1ka0c36khpscnmr8wk7jb-mtcars_tail`. The long cryptographic hash uniquely identifies the build output and is computed based on the content of the derivation and all its inputs. This ensures that the build is fully reproducible.

As a result, building the same derivation on two different machines will yield the same cryptographic hash. You can substitute the built artefact with the derivation that generates it one-to-one, just as in mathematics, where writing $f(2)$ is the same as writing 4 for the function $f(x) := x^2$.

To guarantee that derivations always produce identical outputs, builds must occur in an isolated environment known as a **hermetic sandbox**. This process ensures that the build is unaffected by external factors, such as the state of the host system. This isolation extends to environment variables and even network access. If you need to download data from an API, for example, it will not work from within the build sandbox.

This may seem restrictive, but it makes perfect sense for reproducibility. An API’s output can change over time. For a truly reproducible result, you should obtain the data once, version it, and use that archived data as an input to your analysis.

2.2.4 Other key Nix concepts

We’ve covered derivations, dependencies, closures, the Nix store, and hermetic builds. That’s the core of what makes Nix tick. But there are a few more concepts worth knowing about before we move on:

- Purity: Nix tries very hard to keep builds “pure”: the output only depends on what you explicitly list as inputs. If a build script tries to reach out to the internet or read some random file on your machine, Nix will block it. That can feel restrictive at first, but it’s what guarantees reproducibility.
- Binary caches: You don’t always need to build everything yourself. Think back to the math analogy: if you already know that $f(2) = 4$, there’s no need to compute it again—you can just reuse the result. Nix does the same with binary caches: every build is identified by a unique cryptographic hash of its inputs. That means a prebuilt package fetched from `cache.nixos.org` (or your own cache) is

2 The Nix Package Manager

bit-for-bit identical to what you would have built locally. This is why Nix is both reproducible and fast.

- Garbage collection: Since Nix never overwrites anything in the store, old packages can pile up.
Running `nix-store --gc` will run the garbage collector to free up space.
- Overlays: If you want to tweak a package or add your own without forking all of `nixpkgs`, you can use overlays. They let you extend or override existing definitions in a clean, composable way.
- Flakes: The newer way to define and pin Nix projects. Flakes make it easier to share and reuse Nix setups across machines and repositories.

Together, these features explain why Nix isn't *just another package manager*. It's more like a framework for reproducible environments that can scale from a single project to an entire operating system (called NixOS²).

With these extra concepts in mind, we can now wrap up how Nix ties everything together before moving on to installing it.

2.3 Caveats

While Nix is powerful, there are some limitations and practical hurdles to be aware of if you plan to use it for actual work:

- Hardware acceleration: On non-NixOS systems, it can be difficult to set up GPU acceleration (CUDA, ROCm, OpenCL). Drivers are tightly coupled to the host kernel

²<https://nixos.org/>

and libraries, while Nix builds aim for strict isolation. On NixOS this integration is smoother, but on macOS or Windows you may encounter limitations or extra manual steps.

- macOS-specific issues: Reproducibility is harder to achieve on macOS (both Intel and Apple Silicon) than on Linux. Nix packages often rely on Apple system frameworks (e.g., CoreFoundation, Security) that live outside the Nix store and compromise hermetic builds³. The macOS sandbox is also weaker than Linux's and sometimes leaks system tools like Xcode or Rosetta into builds⁴. Hydra cache coverage is thinner for Darwin platforms, especially Apple Silicon, so cache misses and local builds are much more common⁵. Finally, pinned environments can break after macOS or Xcode updates because of changes in system libraries or compiler flags⁶.

That said, in practice most packages build fine on macOS, and the ecosystem continues to improve. When reproducibility problems do appear, the simplest fix is often just to use another nearby `nixpkgs` revision for pinning that is known to work. This makes the situation less fragile than it might first appear. Another solution would be to use Docker to deploy the Nix environment and use that as a dev container. All of this will be discussed in detail in this book.

- Steep learning curve: The Nix language and ecosystem (flakes, overlays, derivations) can be conceptually difficult

³<https://github.com/NixOS/nixpkgs/issues/67166>

⁴<https://discourse.nixos.org/t/nix-macos-sandbox-issues-in-nix-2-4-and-later/17475>

⁵https://www.reddit.com/r/NixOS/comments/17uxj6q/how_does_nix_package_m/

⁶<https://github.com/NixOS/nix/issues/11679>

2 The Nix Package Manager

if you come from traditional package managers. Even basic customizations require some ramp-up time.

- Disk space and builds: Because Nix never mutates software in place, the store can accumulate large amounts of data. Cache misses sometimes force local builds, which can be slow and resource-intensive. However, it is of course possible to empty the Nix store to recover disk space, and it is also possible to set up your own project cache if you wish so. Setting up your own cache will also be something that we will explore in this book.

These caveats don't diminish Nix's strengths but highlight that its guarantees are strongest on Linux (and thus WSL), and especially on NixOS. On macOS, reproducibility is possible but sometimes requires extra work, a bit of flexibility, and occasionally picking a different `nixpkgs` snapshot.

2.4 In summary

Nix makes it possible to *actually* build software reproducibly. To achieve this, it introduces several core concepts that are quite specific, and definitely worth taking the time to understand.

Nix generates a derivation from a Nix expression through a process called instantiation. During this process, Nix resolves all inputs and computes a unique cryptographic hash from the contents of the derivation and its entire dependency graph. This ensures that even the smallest change results in a distinct derivation.

Once instantiated, the derivation is built in a hermetic environment where only explicitly declared dependencies are available. If a pre-built binary is available through the official Nix cache

(or through another cache you might have added), it gets fetched instead, since building it locally would have resulted in **exactly** the same binary anyways. This makes the build entirely deterministic. After a successful build, Nix stores the output in the Nix store under a unique path determined by its hash. This process is extremely precise: even changing from the comma to the pipe as the separator in a CSV file will result in a different hash for that file, and thus for all the build artifacts that depend on it directly or indirectly.

There are however some caveats, such as hardware acceleration being not necessarily straightforward to configure, or builds that are unfortunately not as reproducible as they could be on macOS. But for each of these issues, there are mitigations.

In the next chapter, we will learn how to install Nix, and use `{rix}` to set up our first reproducible development environments!

The key takeaway is that Nix is a complex tool because it solves a complex problem: ensuring complete reproducibility across different environments and over time. Nix is a demanding mistress, and adding it to your toolbox is not trivial. But the long-term benefits far outweigh the costs. And with my packages `{rix}` and `{rixpress}` (or `ryxpress` for Python), the entry ticket has never been cheaper. You will be in a position to quickly benefit from its power without having to master all its complexities.

3 —

4 Reproducibility with Nix

In this chapter, we will explore the critical need for environment reproducibility in modern workflows. We will learn how to use `{frinx}` to generate `default.nix` files and see how this enables building robust, cross-language environments for data science and software development.

4.1 Why Reproducibility? Why Nix? *(1h30)*

4.1.1 Motivation: Reproducibility in Scientific and Data Workflows

To ensure that a project is reproducible you need to deal with at least four things:

- Make sure that the required/correct version of R (or any other language) is installed;
- Make sure that the required versions of packages are installed;
- Make sure that system dependencies are installed (for example, you'd need a working Java installation to install the `{rJava}` R package on Linux);

4 Reproducibility with Nix

- Make sure that you can install all of this for the hardware you have on hand.

But in practice, one or most of these bullet points are missing from projects. The goal of this course is to learn how to fulfill all the requirements to build reproducible projects.

4.1.2 Problems with Ad-Hoc Tools

Tools like Python’s `venv` or R’s `renv` only deal with some pieces of the reproducibility puzzle. Often, they assume an underlying OS, do not capture system-level dependencies (like `libxml2`, `pandoc`, or `curl`), and require users to “rebuild” their environments from partial metadata. Docker helps but introduces overhead, security challenges, and complexity, and just adding it to your project doesn’t make it reproducible if you don’t explicitly take some precautionary steps.

Traditional approaches fail to capture the entire dependency graph of a project in a deterministic way. This leads to “it works on my machine” syndromes, onboarding delays, and subtle bugs.

4.1.3 Nix, a declarative package manager

Nix is a tool for reproducible builds and development environments, often introduced as a package manager. It captures complete dependency trees, from your programming language interpreter to every system-level library you rely on. With Nix, environments are not recreated from documentation, but rebuilt precisely from code.

4.1 Why Reproducibility? Why Nix? (1h30)

Nix can be installed on Linux distributions, macOS and it even works on Windows if you enable WSL2. In this course, we will use Nix mostly as a package manager (but towards also as a build automation tool).

However Nix has quite a steep learning curve, so this is why for the purposes of this course we are going to use an R package called `{rix}` to set up reproducible environments.

4.1.4 The `rix` package

`{rix}` is an R package (I'm the author) and its goal is to make writing Nix expressions easy. With `{rix}` you can declare the environment you need using the provided `rix()` function, which is the package's main function. Calling it generates a file called `default.nix` which is then used by the Nix package manager to build that environment. Ideally, you would set up such an environment for each of your projects. You can then use this environment to either work interactively, or run R or Python scripts. It is possible to have as many environments as projects, and software that is common to environments will simply be re-used and not get re-installed to save space. Environments are isolated from each other, but can still interact with your system's files, unlike with Docker where a volume must be mounted. While this is useful, it can sometimes lead to issues. For example, if you already have R installed, and a user library of R packages, more caution is required to properly use environments managed by Nix.

You don't need to have R installed or be an R user to use `{rix}`. If you have Nix installed on your system, it is possible to "drop" into a temporary environment with R and `{rix}` available and generate the required Nix expression from there.

4 Reproducibility with Nix

But first, let's install Nix and try to use temporary shells.

4.1.5 Installing Nix

4.1.5.1 For Windows users only: some prerequisites

If you are on Windows, you need the Windows Subsystem for Linux 2 (WSL2) to run Nix. If you are on a recent version of Windows 10 or 11, you can simply run this as an administrator in PowerShell:

```
wsl --install
```

You can find further installation notes at this official MS documentation.

I recommend to activate `systemd` in Ubuntu WSL2, mainly because this supports other users than `root` running Nix. To set this up, please do as outlined this official Ubuntu blog entry:

```
# in WSL2 Ubuntu shell

sudo -i
nano /etc/wsl.conf
```

This will open the `/etc/wsl.conf` in a nano, a command line text editor. Add the following line:

```
[boot]
systemd=true
```

4.1 Why Reproducibility? Why Nix? (1h30)

Save the file with CTRL-O and then quit nano with CTRL-X. Then, type the following line in powershell:

```
wsl --shutdown
```

and then relaunch WSL (Ubuntu) from the start menu. For those of you running Windows, we will be working exclusively from WSL2 now. If that is not an option, then I highly recommend you set up a virtual machine with Ubuntu using VirtualBox for example, or dual-boot Ubuntu.

Installing (and uninstalling) Nix is quite simple, thanks to the installer from Determinate Systems, a company that provides services and tools built on Nix, and works the same way on Linux (native or WSL2) and macOS.

4.1.5.2 Actually installing Nix

Do not use your operating system's package manager to install Nix. Instead, simply open a terminal and run the following line (on Windows, run this inside WSL):

```
curl --proto '=https' --tlsv1.2 -sSf \
-L https://install.determinate.systems/nix | \
sh -s -- install
```

Then, install the `cachix` client and configure the `rstats-on-nix` cache: this will install binary versions of many R packages which will speed up the building process of environments:

4 Reproducibility with Nix

```
nix-env -iA cachix -f
↳ https://cachix.org/api/v1/install
```

then use the cache:

```
cachix use rstats-on-nix
```

You only need to do this once per machine you want to use {rix} on. Many thanks to Cachix for sponsoring the `rstats-on-nix` cache!

If you get this warning when trying to install software with Nix:

```
warning: ignoring the client-specified setting
'trusted-public-keys', because it is a restricted
setting and you are not a trusted user
warning: ignoring untrusted substituter
'https://rstats-on-nix.cachix.org', you are not a
trusted user.
Run `man nix.conf` for more information on the
`substituters` configuration option.
warning: ignoring the client-specified setting
'trusted-public-keys', because it is a restricted
setting and you are not a trusted user
```

Then this means that configuration was not successful. You need to add your user to `/etc/nix/nix.custom.conf`:

```
sudo vim /etc/nix/nix.custom.conf
```

then simply add this line in the file:

```
trusted-users = root YOURUSERNAME
```

where `YOURUSERNAME` is your current login user name.

4.1.6 Temporary shells

You now have Nix installed; before continuing, let's see if everything works (close all your terminals and reopen them) by dropping into a temporary shell with a tool you likely have not installed on your machine.

Open a terminal and run:

```
which sl
```

you will likely see something like this:

```
which: no sl in ....
```

now run this:

```
nix-shell -p sl
```

and then again:

```
which sl
```

this time you should see something like:

```
/nix/store/cndqpx74312xkrrgp842ifinkd4cg89g-sl-5.05/bin/sl
```

This is the path to the `sl` binary installed through Nix. The path starts with `/nix/store`: the *Nix store* is where all the software installed through Nix is stored. Now type `sl` and see what happens!

You can find the list of available packages here.

4.2 Session 1.2 – Dev Environments with Nix (1h30)

4.2.1 Some Nix concepts

While temporary shells are useful for quick testing, this is not how Nix is typically used in practice. Nix is a declarative package manager: users specify what they want to build, and Nix takes care of the rest.

To do so, users write files called `default.nix` that contain the a so-called Nix expression. This expression will contain the definition of a (or several) *derivations*.

In Nix terminology, a derivation is *a specification for running an executable on precisely defined input files to repeatably produce output files at uniquely determined file system paths.* (source)

In simpler terms, a derivation is a recipe with precisely defined inputs, steps, and a fixed output. This means that given identical inputs and build steps, the exact same output will always be produced. To achieve this level of reproducibility, several important measures must be taken:

- All inputs to a derivation must be explicitly declared.
- Inputs include not just data files, but also software dependencies, configuration flags, and environment variables, essentially anything necessary for the build process.
- The build process takes place in a *hermetic* sandbox to ensure the exact same output is always produced.

The next sections of this document explain these three points in more detail.

4.2.2 Derivations

Here is an example of a *simple* Nix expression:

```
let

pkgs = import (fetchTarball
    "https://github.com/rstats-on-nix/nixpkgs/archive/2025-0
    {});

in

pkgs.stdenv.mkDerivation {
    name = "filtered_mtcars";
    buildInputs = [ pkgs.gawk ];
    dontUnpack = true;
    src = ./mtcars.csv;
    installPhase = ''
        mkdir -p $out
        awk -F',' 'NR==1 || $9=="1" { print }' $src >
    $out/filtered.csv
    '';
}
}
```

I won't go into details here, but what's important is that this code uses `awk`, a common Unix data processing tool, to filter the `mtcars.csv` file to keep only rows where the 9th column (the `am` column) equals 1. As you can see, a significant amount of boilerplate code is required to perform this simple operation. However, this approach is completely reproducible: the dependencies are declared and pinned to a specific dated branch of our `rstats-on-nix/nixpkgs` fork (more on this later), and the only thing that could make this pipeline fail (though it's a bit of

4 Reproducibility with Nix

a stretch to call this a *pipeline*) is if the `mtcars.csv` file is not provided to it. This expression can be *instantiated* into a derivation, and the derivation is then built into the actual output that interests us, namely the filtered `mtcars` data.

The derivation above uses the `Nix` builtin function `mkDerivation`: as its name implies, this function *makes a derivation*. But there is also `mkShell`, which is the function that builds a shell instead. `Nix` expressions that built a shell is the kind of expressions `{rix}` generates for you.

4.2.3 Using `{rix}` to generate development environments

If you have successfully installed `Nix`, but don't have yet `R` installed on your system, you could install `R` as you would usually do on your operating system, and then install the `{rix}` package, and from there, generate project-specific expressions and build them. But you could also install `R` using `Nix`. Actually, I would even recommend you uninstall `R` and delete all your packages from your computer and only manager `R` environments using `Nix`.

Running the following line in a terminal will drop you in an interactive `R` session that you can use to start generating expressions:

```
nix-shell -p R rPackages.rix
```

This will drop you in a temporary shell with `R` and `{rix}` available. Navigate to an empty directory to help a project, call it `rix-session-1`:

```
mkdir rix-session-1
```

and start R and load {rix}:

R

```
library(rix)
```

you can now generate an expression by running the following code:

```
rix(
  date = "2025-08-04",
  r_pkgs = c("dplyr", "ggplot2"), # add
  ↵ languageserver if you plan to use VS Code
  py_conf = list(
    py_version = "3.13",
    py_pkgs = c("polars", "great-tables")
  ),
  ide = "none",
  project_path = ".",
  overwrite = TRUE
)
```

This will write a file called `default.nix` in your project's directory. This `default.nix` contains a Nix expression which will build a shell that comes with R, `{dplyr}` and `{ggplot2}` as they were on the the 4th of August 2025 on CRAN. This will also add Python 3.13 and the `polars` and `great-tables` Python packages as they were at the time in `nixpkgs` (more

4 Reproducibility with Nix

on this later). Finally, the `ide` argument is set to "none", because we don't want to use Nix to manage an IDE. However, you should know that this is possible and useful in some cases. See this vignette for learning how to setup your IDE with Nix if you wish to do so. At the end of this chapter we will learn how to set up Positron.

4.2.4 Using `nix-shell` to Launch Environments

Once your file is in place, simply run:

```
nix-shell
```

This gives you an isolated shell session with all declared packages available. You can test code, explore APIs, or install further tools within this session.

To remove the packages that were installed, call `nix-store --gc`. This will call the garbage collector. If you want to avoid that an environment gets garbage-collected, use `nix-build` instead of `nix-shell`. This will create a symlink called `result` in your project's root directory and `nix-store --gc` won't garbage-collect this environment until you manually remove `result`.

4.2.5 Pinning with `nixpkgs`

To ensure long-term reproducibility, a pinned the version of Nixpkgs is used:

```
let
  pkgs = import (fetchTarball
    ↵  "https://github.com/rstats-on-nix/nixpkgs/archive/2025-
    ↵  {});
in
...
```

This is done automatically by `{rix}`. You could change the date manually if you prefer, but I would recommend to always regenerate the `default.nix` using `{rix}`.

4.2.6 Installing Python packages not available via nixpkgs (impure)

Not all Python packages can be installed through Nix; unlike CRAN, Pypi doesn't get automatically mirrored and individual packages fixed by volunteers. Instead, specific Python packages get packaged individually for Nix. Thus, it could very well be the case that a specific Python package (or version of a Python package) that you need for a project is not available via `nixpkgs`.

In this case, it is still possible to use Python-specific package managers, like `uv`, to install packages. This is also useful if you work on a project with colleagues that use `uv` and that don't want (yet) to use Nix. `uv`, is 10-100x faster than `pip` and also generates a lock file for improved reproducibility.

The idea is to install `uv` in your shell (but not any Python nor Python packages):

4 Reproducibility with Nix

```
rix(
  ...
  system_pkgs = c("uv"),
  ...
)
```

And then use uv from your shell as you would usually. We recommend specifying Python packages in a `requirements.txt` file, and specifying explicit versions (e.g., `scipy==1.11.4`). Finally, we also recommend setting a shell hook to set up the virtual environment and install the packages from the `requirements.txt` when entering the shell (mind the quotes):

```
rix(
  ...
  system_pkgs = c("uv"),
  shell_hook = "
    if [ ! -f pyproject.toml ]; then
      uv init --python 3.13.5 # or whichever
    ↵ Python version you need
    fi
    uv add --requirements requirements.txt
    # Create alias so python uses uvs environment
    alias python='uv run python'
  ",
)
```

After running `nix-shell`, uv should initialize a Python project with the specified Python version and install the packages listed in `requirements.txt` within the nix environment. This will take place each time `nix-shell` is called, however this will be cached and not installed each time.

4.2 Session 1.2 – Dev Environments with Nix (1h30)

To make sure everything works fine, you could simply start a Python interpreter and try to load `numpy`. This should work fine, but if it doesn't, the following error could be raised:

`ImportError:`

IMPORTANT: PLEASE READ THIS FOR ADVICE ON HOW TO
↪ SOLVE THIS ISSUE!

Importing the `numpy` C-extensions failed. This error
↪ can happen for
many reasons, often due to issues with your setup or
↪ how NumPy was
installed.

We have compiled some common reasons and
↪ troubleshooting tips at:

↪ <https://numpy.org/devdocs/user/troubleshooting-importerror>

Please note and check the following:

- * The Python version is: `Python3.13` from
↪ `"/home/user/projects/rix_uv/.venv/bin/python3"`
- * The NumPy version is: `"2.2.6"`

and make sure that they are the versions you expect.
↪ Please carefully study the
documentation linked above for further help.

Original error was: `libstdc++.so.6: cannot open`
↪ `shared object file: No such file or directory`

4 Reproducibility with Nix

That is an issue when using *wheels* (wheels are binaries of Python packages that get installed by default using uv). These wheels expect certain libraries to be in certain places. One way to solve this is to add the following to your shell hook:

```
shellHook = ''  
  # Export LD_LIBRARY is required for python  
  #   packages that dynamically load libraries,  
  #   such as numpy  
  export  
  ↵ LD_LIBRARY_PATH="${pkgs.lib.makeLibraryPath  
  ↵ (with pkgs; [ zlib gcc.cc glibc stdenv.cc.cc  
  ↵ ])}":LD_LIBRARY_PATH;  
  ...  
'';
```

Your environment should now work.

If this seems complicated: yes, and that is actually exactly the type of problems that Nix aims to solve. However, there are just too many Python packages to automate their inclusion into nixpkgs like how it's done for R. If you can, prefer using the Python packages included in nixpkgs.

4.3 Configuring your IDE

4.3.1 Pre-requisites

We now need to configure an IDE to use both our Nix shells as development environments, and GitHub Copilot. You are free to use whatever IDE you want but the instructions below are going

to focus on Positron, which is a fork of VS Code geared towards data science. It works well with both Python and R and makes it quite easy to choose the right R or Python interpreter (which you'll have to do to make sure you're using the one provided by Nix, see here).

If you want to use VS Code proper, you can follow all the instructions here, but you need to install the REditorSupport and the Python extension. You will also need to add the `{languageserver}` R package to the shell:

```
rix(
  date = "2025-08-04",
  r_pkgs = c("dplyr", "ggplot2", "languageserver"),
  py_conf = list(
    py_version = "3.13",
    py_pkgs = c("polars", "great-tables")
  ),
  ide = "none",
  project_path = ".",
  overwrite = TRUE
)
```

and re-create the `default.nix` file.

You'll need to carefully select the Python interpereter provided by Nix, and not the other environments that you might already have on your system. Also, we recommend you uninstall R if it's installed system-wide and also remove your local library of packages and instead only use dedicated Nix shells to manage your projects. While we made our possible for Nix shells to not interfere with a system-installed R, we recommend users go into the habit of taking some minutes at the start of a project to properly set up their development environment.

4 Reproducibility with Nix

On Windows, you need to install Positron on Windows, not inside WSL. Then, install the Open Remote extension right from Positron itself.

4.3.2 direnv

Once Positron is installed, you need to install a piece of software called `direnv`: `direnv` will automatically load Nix shells when you open a project that contains a `default.nix` file in an editor. It works on any operating system and many editors support it, including Positron. If you're using Windows, install `direnv` in WSL (even though you've just installed Positron for Windows). To install `direnv` run this command:

```
nix-env -f '<nixpkgs>' -iA direnv
```

This will install `direnv` and make it available even outside of Nix shells!

Then, we highly recommend to install the `nix-direnv` extension:

```
nix-env -f '<nixpkgs>' -iA nix-direnv
```

It is not mandatory to use `nix-direnv` if you already have `direnv`, but it'll make loading environments much faster and seamless.

Finally, if you haven't used `direnv` before, don't forget this last step to make your terminal detect and load `direnv` automatically.

Then, in Positron, install the direnv extension. Finally, add a file called `.envrc` and simply write the following two lines in it (this `.envrc` file should be in the same folder as your project's `default.nix`):

```
use nix
mkdir $TMP
```

in it. On Windows, *remotely connect to WSL* first, but on other operating systems, simply open the project's folder using **File > Open Folder...** and you will see a pop-up stating `direnv: /PATH/TO/PROJECT/.envrc is blocked` and a button to allow it. Click **Allow** and then open an R script. You might get another pop-up asking you to restart the extension, so click **Restart**. Be aware that at this point, `direnv` will run `nix-shell` and so will start building the environment. If that particular environment hasn't been built and cached yet, it might take some time before Code will be able to interact with it. You might get yet another popup, this time from the R Code extension complaining that R can't be found. In this case, simply restart Code and open the project folder again: now it should work every time.

Sometimes, depending on platform and installed packages, Positron might not be able to correctly load a Nix-managed Python interpreter. This happens when it tries to load its own version of `ipykernal`. To make sure this doesn't cause issues, add these lines to the shell's definition in `default.nix`:

```
shellHook = ''
# Export LD_LIBRARY is required for python
↳ packages that dynamically load libraries,
↳ such as numpy
```

4 Reproducibility with Nix

```
export
↳ LD_LIBRARY_PATH="${pkgs.lib.makeLibraryPath
↳ (with pkgs; [ zlib gcc.cc glibc stdenv.cc.cc
↳ ])}"':LD_LIBRARY_PATH;
...
'';
```

4.3.3 In summary and next steps

For a new project, simply repeat this process:

- Generate the project's `default.nix` file;
- Build it using `nix-build`;
- Create an `.envrc` and write the two lines from above in it;
- Open the project's folder in Code and click allow when prompted;
- Restart the extension and Code if necessary.

Another option is to create the `.envrc` file and write `use nix` in it, then open a terminal, navigate to the project's folder, and run `direnv allow`. Doing this before opening Positron should not prompt you anymore.

If you're on Windows, using Positron like this is particularly interesting, because it allows you to install Positron on Windows as usual, and then you can configure it to interact with a Nix shell, even if it's running from WSL. This is a very seamless experience.

Now configure Positron to use GitHub Copilot: [click here](#).

You can enable LLM tab-completion if you want. I'm not a fan of this, I prefer simply using the interface and using LLMs as an advanced rubber duck: <https://github.com/copilot/>.

4.4 Starting a new project:

To start a new project, create a new directory, and `cd` into it. Then, in that folder, create a new script called `gen-env.R` and add the following:

```
library(rix)

rix(
  date = "2025-09-22",
  r_pkgs = c(
    "languageserver", # if you want to use VS Code
    "tidyverse" # or whatever packages
  ),
  py_conf = list(
    py_version = "3.13",
    py_pkgs = c("polars", "great-tables")
  ),
  ide = "none",
  project_path = ".",
  overwrite = TRUE,
  print = TRUE
)
```

Then call:

4 Reproducibility with Nix

```
nix-shell -I
↳ nixpkgs=https://github.com/rstats-on-nix/nixpkgs/archive/re
↳ -p R rPackages.rix
```

providing a url to my fork of `nixpkgs` makes sure that you get the latest `rix` version (change the date in the url to a more recent one if needed). Start R by typing `R` in the shell, and then `source("gen-env.R")`. That will create a `default.nix` which you can then drop into using `nix-shell`. Then create an `.envrc` and add the lines from before, and you're done!

Note: if you want to use `uv` instead of Nix to manage the Python packages, follow the instructions above.

4.5 Hands-On Exercises

1. Start a temporary shell with R and `{rix}` again using `nix-shell -p R rPackages.rix`. Start an R session (by typing `R`) and then load the `{rix}` package (using `library(rix)`). Run the `available_dates()` function: using the latest available date, generate a new `default.nix`.
2. Inside of an activated shell, type `which R` and `echo $PATH`. Explore what is being added to your environment. What is the significance of paths like `/nix/store/...`?
3. Break it on purpose: generate a new environment with a wrong R package name, for example `dplyrnaught`. Try to build the environment. What happens?
4. Go to <https://search.nixos.org/packages> and look for packages that you usually use for your projects to see if they are available.

5 Reproducible Development Environments with rix

5.1 Introduction

5.1.1 The Reproducibility Challenge

Reproducibility in research and data science exists on a continuum. At one end, authors might only describe their methods in prose. Moving along the spectrum, they might share code, then data, and finally what we call a *computational environment*: the complete set of software required to execute an analysis.

Even when researchers share code and data, they rarely specify the full software stack: the exact version of R, all package versions, and crucially, the system-level dependencies. Yet differences in any of these can lead to divergent results from the same code.

Tools like `{renv}` address part of the problem—they capture R package versions in a lockfile. But `{renv}` doesn’t manage the R version itself (you need `rig` for that), and neither handles system libraries. If `{sf}` requires GDAL 3.0 but your system has 2.4, `{renv}` can’t help. And if your project uses both R and

Python? Now you’re coordinating multiple package managers, each with its own configuration.

5.1.2 Component Closures: The Nix Approach

This is where Nix shines. Nix deploys *component closures*: when you install a package, Nix also installs all its dependencies, their dependencies, and so on. Think of it like packing for a trip—traditional package managers assume you’ll find essentials at your destination, while Nix packs everything you need.

As the original Nix paper explains:

The idea is to always deploy component closures: if we deploy a component, then we must also deploy its dependencies, their dependencies, and so on. Since closures are self-contained, they are the units of complete software deployment.

This means when you install `{sf}` through Nix, you automatically get the correct versions of GDAL, GEOS, and PROJ—no manual system configuration needed.

5.1.3 The Polyglot Challenge

Modern data science is increasingly polyglot. Research shows that data scientists use, on average, nearly two programming languages in their work, with R and Python being the most common combination. Python dominates machine learning, R excels at statistical modelling, and Julia offers high-performance numerics. Projects increasingly combine these strengths.

This creates a reproducibility challenge: a project using R, Python, and Quarto requires coordinating multiple package managers. Nix solves this by providing a unified framework for all languages and system tools.

5.1.4 Enter `rix`

However, Nix has a steep learning curve. Its functional programming language can be daunting for researchers focused on their analysis, not system administration.

`{rix}` bridges this gap. It's an R package that generates Nix expressions from intuitive R function calls. You describe *what* you want, and `{rix}` figures out *how* to express it in Nix. The workflow is simple:

1. **Declare** your environment using `rix()`
2. **Build** the environment using `nix-build`
3. **Use** the environment using `nix-shell`

This chapter covers everything you need to know to create project-specific, reproducible development environments for your R projects.

5.2 The `rix()` Function

The `rix()` function is the heart of the package. It generates a `default.nix` file—a Nix expression that defines your development environment. Here are its main arguments:

- `r_ver`: the version of R you need (or use `date` instead)
- `date`: a date corresponding to a CRAN snapshot

5 Reproducible Development Environments with *rix*

- `r_pkgs`: R packages to install from CRAN/Bioconductor
- `system_pkgs`: system tools like `quarto` or `git`
- `git_pkgs`: R packages to install from GitHub
- `local_r_pkgs`: local `.tar.gz` packages to install
- `tex_pkgs`: TeXLive packages for literate programming
- `ide`: which IDE to configure ("`rstudio`", "`code`", "`positron`", "`none`")
- `project_path`: where to save the `default.nix` file
- `overwrite`: whether to overwrite an existing `default.nix`

Let's create our first environment. Suppose you need R with `{dplyr}` and `{ggplot2}`:

```
library(rix)

rix(
  r_ver = "4.4.2",
  r_pkgs = c("dplyr", "ggplot2"),
  ide = "code",
  project_path = ".",
  overwrite = TRUE
)
```

This generates two files:

1. `default.nix`: The Nix expression defining your environment
2. `.Rprofile`: Created by `rix_init()` (called automatically), this file prevents conflicts with any system-installed R packages

The `.Rprofile` is important: it ensures that packages from your user library don't get loaded into the Nix environment, and it redefines `install.packages()` to throw an error—because you should never install packages that way in a Nix environment.

5.3 Choosing an R Version or Date

The `r_ver` argument (or alternatively `date`) controls which version of R and which package versions you'll get. Here's a summary of the options:

<code>r_ver / date</code>	Intended Use	R Version	Package Versions
"latest-upstream"	New project, versions don't matter	Current/previous	Up to 6 months old
"4.4.2" (or similar)	Reproduce old project or start new	Specified version	Up to 2 months old
<code>date = "2024-12-14"</code>	Precise CRAN snapshot	Current at that date	Exact versions from date
"bleeding-edge"	Develop against latest CRAN	Always current	Always current
"frozen-edge"	Latest CRAN, manual updates	Current at generation	Current at generation
"r-devel"	Test against R development version	R-devel	Always current

To see which R versions are available:

```
available_r()
```

To see which dates are available for snapshotting:

```
available_dates()
```

5.3.1 Using dates vs versions

Using a specific date is often the best choice for reproducibility. When you specify a date, you get the exact state of CRAN on that day:

```
rix(  
  date = "2024-12-14",  
  r_pkgs = c("dplyr", "ggplot2"),  
  ide = "code",  
  project_path = ".",  
  overwrite = TRUE  
)
```

5.3.2 The rstats-on-nix fork

When you use a specific R version or date (rather than "latest-upstream"), {rix} uses our rstats-on-nix fork of nixpkgs rather than the upstream repository. This fork:

- Provides better compatibility for older packages
- Includes many fixes, especially for Apple Silicon
- Offers newer R releases faster than the official channels
- Snapshots CRAN more frequently

5.4 Installing R Packages

5.4.1 From CRAN/Bioconductor

The simplest case—just list package names in `r_pkgs`:

```
rix(
  r_ver = "4.4.2",
  r_pkgs = c("dplyr", "ggplot2", "tidyverse", "readr"),
  ide = "code",
  project_path = "."
)
```

Both CRAN and Bioconductor packages can be specified this way.

5.4.2 Installing archived versions

Need a specific old version of a package? Use the `@` syntax:

```
rix(
  r_ver = "4.2.1",
  r_pkgs = c("dplyr@0.8.0", "janitor@1.0.0"),
  ide = "none",
  project_path = "."
)
```

This will install `{dplyr}` version 0.8.0 from the CRAN archives. Note that archived packages are built from source, which may fail for packages requiring compilation.

5.4.3 Installing from GitHub

For packages on GitHub, use the `git_pkgs` argument with a list containing the package name, repository URL, and commit hash:

```
rix(
  r_ver = "4.4.2",
  r_pkgs = c("dplyr", "ggplot2"),
  git_pkgs = list(
    list(
      package_name = "housing",
      repo_url =
        "https://github.com/rap4all/housing/",
      commit =
        "1c860959310b80e67c41f7bbdc3e84cef00df18e"
    )
  ),
  ide = "none",
  project_path = "."
)
```

Always specify a commit hash, not a branch name. This ensures reproducibility—branch names can change, but commits are immutable.

If the R package lives in a subfolder of the repository, append the subfolder to the URL:

```
git_pkgs = list(
  package_name = "BPCells",
  repo_url = "https://github.com/bnprks/BPCells/r",
  # Note the /r suffix
```

```

commit =
  ↴ "16faeade0a26b392637217b0caf5d7017c5bdf9b"
)

```

5.4.4 Installing local packages

For local `.tar.gz` archives, place them in the same directory as your `default.nix` and use `local_r_pkgs`:

```

nix(
  r_ver = "4.3.1",
  local_r_pkgs = c("mypackage_1.0.0.tar.gz"),
  ide = "none",
  project_path = "."
)

```

5.4.5 Why NOT to use `install.packages()`

It's crucial to understand: **never call `install.packages()` from within a Nix environment.** Here's why:

- 1. Declarative environments:** If you install packages imperatively, your `default.nix` no longer matches your actual environment.
- 2. Leaking packages:** Packages installed via `install.packages()` go to your user library, not the Nix environment. They'll be visible to *all* Nix shells, breaking isolation.
- 3. Reproducibility:** The whole point of Nix is that your environment is fully defined by the `default.nix`. Ad-hoc installations defeat this.

Instead, add packages to your `rix()` call and rebuild the environment.

5.5 System Tools and TexLive

5.5.1 Adding system packages

Need command-line tools like Quarto or Git? Add them via `system_pkgs`:

```
rix(  
  r_ver = "latest-upstream",  
  r_pkgs = c("quarto"),  
  system_pkgs = c("quarto", "git"),  
  ide = "none",  
  project_path = ".")
```

Note that here we install both the R `{quarto}` package and the `quarto` command-line tool—they’re different things!

To find available packages, search at search.nixos.org.

5.5.2 TexLive for literate programming

For PDF output from Quarto, R Markdown, or Sweave, you need TexLive packages:

```
rix(
  r_ver = "latest-upstream",
  r_pkgs = c("quarto"),
  system_pkgs = "quarto",
  tex_pkgs = c("amsmath", "framed", "fvextra"),
  ide = "none",
  project_path = "."
)
```

This installs the `scheme-small` TexLive distribution plus the specified packages.

5.5.3 Python and Julia integration

`{rix}` can also add Python or Julia to your environment:

```
rix(
  date = "2025-02-17",
  r_pkgs = "ggplot2",
  py_conf = list(
    py_version = "3.12",
    py_pkgs = c("polars", "great-tables")
  ),
  ide = "positron",
  project_path = "."
)
```

This creates a polyglot environment with R, Python, and the specified packages for each.

5.5.4 Installing Python packages with **uv** (impure)

Not all Python packages are available through Nix. Unlike CRAN, PyPI doesn't get automatically mirrored—individual packages must be packaged by volunteers. If a Python package you need isn't in `nixpkgs`, you can use `uv` as an escape hatch.

This approach is also useful when collaborating with colleagues who use `uv` but haven't adopted Nix yet. `uv` is 10–100x faster than `pip` and generates a lock file for improved reproducibility.

The idea is to install `uv` in your shell (but not Python or Python packages through Nix):

```
rix(  
    date = "2025-02-17",  
    r_pkgs = "ggplot2",  
    system_pkgs = c("uv"),  
    ide = "none",  
    project_path = ".")
```

Then use `uv` from within your shell. We recommend:

1. Specify Python packages in a `requirements.txt` file with explicit versions (e.g., `scipy==1.11.4`)
2. Set up a shell hook to automatically configure the virtual environment

Here's a complete example with a shell hook:

```

rix(
    date = "2025-02-17",
    r_pkgs = "ggplot2",
    system_pkgs = c("uv"),
    shell_hook =
        if [ ! -f pyproject.toml ]; then
            uv init --python 3.13.5
        fi
        uv add --requirements requirements.txt
        alias python='uv run python'
    ",
    ide = "none",
    project_path = "."
)

```

After running `nix-shell`, `uv` initialises a Python project with the specified version and installs packages from `requirements.txt`. This happens each time you enter the shell, but `uv` caches everything so it's nearly instant after the first run.

5.5.4.1 Troubleshooting wheel issues

When using wheels (pre-compiled Python packages), you may encounter errors like:

```
ImportError: libstdc++.so.6: cannot open shared
object file
```

This happens because wheels expect certain libraries in certain locations. Add this to your shell hook to fix it:

```
shellHook = ''  
  export LD_LIBRARY_PATH="${pkgs.lib.makeLibraryPath  
  ↵  (with pkgs; [  
    zlib gcc.cc glibc stdenv.cc.cc  
  ])}"${LD_LIBRARY_PATH  
  # ... rest of your hook  
'';
```

If this seems complicated: yes, it is. This is exactly the kind of problem Nix aims to solve. When possible, prefer Python packages included in `nixpkgs`.

5.6 Building and Using Environments

5.6.1 Building with `nix-build`

Once you have a `default.nix`, build the environment:

```
nix-build
```

This downloads/builds all required packages and creates a `result` symlink in your project directory. The `result` file prevents the environment from being garbage-collected.

5.6.2 Entering with `nix-shell`

To use the environment interactively:

```
nix-shell
```

You'll drop into a shell where R and all your packages are available. Type R to start an R session, or start your editor directly.

5.6.3 Running scripts

You can run scripts directly without entering the shell:

```
nix-shell default.nix --run "Rscript analysis.R"
```

Or run a {targets} pipeline:

```
nix-shell default.nix --run "Rscript -e  
  ↵  'targets::tar_make()'"
```

5.6.4 Pure shells for complete isolation

By default, nix-shell can still see programs installed on your system. For complete isolation:

```
nix-shell --pure
```

This hides everything not explicitly included in your environment.

5.6.5 Garbage collection

Nix never deletes old packages automatically. To clean up:

1. Delete the `result` symlink
2. Run `nix-store --gc`

This removes all packages that are no longer referenced by any environment.

5.7 Converting `renv` Projects

If you have existing projects using `{renv}`, the `renv2nix()` function can help you migrate:

```
renv2nix(  
    renv_lock_path = "path/to/project/renv.lock",  
    project_path = "path/to/new_nix_project"  
)
```

5.7.1 Recommended workflow

1. Copy the `renv.lock` file to a new, empty folder
2. Run `renv2nix()` pointing to that folder
3. Build the environment with `nix-build`

Do **not** convert in the same folder as the original `{renv}` project—the generated `.Rprofile` will conflict with `{renv}`'s `.Rprofile`.

5.7.2 Caveats

- Package versions may not match exactly due to how Nix handles snapshotting
- If the `renv.lock` lists an old R version but recent packages, use the `override_r_ver` argument to specify a more appropriate R version

5.8 Summary

In this chapter, we learned how to use `{rix}` to create reproducible development environments:

- The `rix()` function generates `default.nix` files from R
- Choose R versions via `r_ver` or precise CRAN snapshots via `date`
- Install packages from CRAN, Bioconductor, GitHub, or local archives
- Add system tools, TexLive packages, and even Python/Julia
- Build with `nix-build`, use with `nix-shell`
- Convert existing `{renv}` projects with `renv2nix()`

The key insight is that your environment is now *declared* in a text file. Anyone with that file can recreate the exact same environment, on any machine, at any point in the future. That's the power of Nix.

6 Functional Programming

In this chapter, we will see why functional programming is crucial for reproducible, testable, and collaborative data science. We will compare how to write self-contained, “pure” functions in both R and Python, and how to use functional concepts like `map`, `filter`, and `reduce` to replace error-prone loops. Finally, we will discuss how writing functions makes your code easier to review, debug, and even generate with LLMs.

6.1 Introduction: From Scripts to Functions

In the previous chapter, we learned how to create reproducible development environments with `{rix}`. We can now ensure everyone has the *exact same tools*—R, Python, system libraries—to run our code.

But having the right tools is only half the battle. Now we turn to **writing reproducible code itself**. A common way to start a data analysis is by writing a script: a sequence of commands executed from top to bottom.

```
# R script example  
library(dplyr)
```

6 Functional Programming

```
data(mtcars)
heavy_cars <- filter(mtcars, wt > 4)
mean_mpg_heavy <- mean(heavy_cars$mpg)
print(mean_mpg_heavy)
```

```
# Python script example
import pandas as pd
mtcars = pd.read_csv("mtcars.csv")
heavy_cars = mtcars[mtcars['wt'] > 4]
mean_mpg_heavy = heavy_cars['mpg'].mean()
print(mean_mpg_heavy)
```

This works, but it has a hidden, dangerous property: `state`. The script relies on variables like `heavy_cars` existing in the environment, making the code hard to reason about, debug, and test.

6.1.1 The Notebook Problem

If scripting with state is a crack in the foundation of reproducibility, then using computational notebooks is a gaping hole.

Notebooks like Jupyter introduce an even more insidious form of state: the cell execution order. You can execute cells out of order, meaning the visual layout of your code has no relation to how it actually ran. This is a recipe for non-reproducible results and a primary cause of the “it worked yesterday, why is it broken today?” problem.

6.1.2 The Functional Solution

The solution is to embrace a paradigm that minimises state: **Functional Programming (FP)**. Instead of a linear script, we structure our code as a collection of self-contained, predictable functions.

The power of FP comes from the concept of **purity**, borrowed from mathematics. A mathematical function has a beautiful property: for a given input, it always returns the same output. `sqrt(4)` is always 2. Its result doesn't depend on what you calculated before or on a random internet connection.

Our Nix environments handle the “right library” problem; purity handles the “right logic” problem. Our goal is to write our analysis code with this same level of rock-solid predictability.

6.1.3 FP vs OOP: Transformations vs Actors

To appreciate what FP brings, it helps to contrast it with **Object-Oriented Programming (OOP)**, arguably the dominant paradigm in many software systems.

OOP organises computation around *who does what*—a network of objects communicating with each other and managing their own internal state. You send a message to an object, asking it to perform an action, without needing to know how it works internally.

Functional programming, by contrast, organises computation around *how data changes*. It replaces a network of interacting objects with a flow of transformations: data goes in, data comes out, and nothing else changes in the process.

This shift is especially powerful in data science:

6 Functional Programming

- Analyses are naturally expressed as **pipelines of transformations** (cleaning, filtering, aggregating, modelling)
- Pure functions make results **reproducible**—same inputs always yield same outputs
- **Immutability** prevents accidental side effects on shared data
- Because transformations can be composed, tested, and reused independently, FP encourages **modular, maintainable** analysis code

6.1.4 Why Does This Matter for Data Science?

Adopting a functional style brings massive benefits:

1. **Unit Testing is Now Possible:** You can't easily test a 200-line script. But you *can* easily test a small function that does one thing.
2. **Code Review is Easier:** A Pull Request that just adds or modifies a single function is simple for your collaborators to understand and approve.
3. **Working with LLMs is More Effective:** It's incredibly effective to ask, “Write a Python function that takes a pandas DataFrame and a column name, and returns the mean of that column, handling missing values. Also, write three pytest unit tests for it.”
4. **Readability:** Well-named functions are self-documenting:

```
starwars %>%
  group_by(species) %>%
  summarize(mean_height = mean(height))
```

is instantly understandable. The equivalent `for` loop is a puzzle.

6.2 Purity and Side Effects

A **pure function** has two rules:

1. It only depends on its inputs. It doesn't use any "global" variables defined outside the function.
2. It doesn't change anything outside of its own scope. It doesn't modify a global variable or write a file to disk. This is called having "no side effects."

Consider this "impure" function in Python:

```
# IMPURE: Relies on a global variable
discount_rate = 0.10

def calculate_discounted_price(price):
    return price * (1 - discount_rate) # What if
    ↵ discount_rate changes?

print(calculate_discounted_price(100))
# > 90.0
discount_rate = 0.20 # Someone changes the state
print(calculate_discounted_price(100))
# > 80.0 -- Same input, different output!
```

The pure version passes *all* its dependencies as arguments:

6 Functional Programming

```
# PURE: All inputs are explicit arguments
def calculate_discounted_price_pure(price, rate):
    return price * (1 - rate)

print(calculate_discounted_price_pure(100, 0.10))
# > 90.0
print(calculate_discounted_price_pure(100, 0.20))
# > 80.0
```

Now the function is predictable and self-contained.

6.2.1 Handling “Impure” Operations like Randomness

Some operations, like generating random numbers, are inherently impure. Each time you run `rnorm(10)` or `numpy.random.rand(10)`, you get a different result.

The functional approach is not to avoid this, but to *control* it by making the source of impurity (the random seed) an explicit input.

In R, the `{withr}` package helps create a temporary, controlled context:

```
library(withr)

# This function is now pure! For a given seed, the
# output is always the same.
pure_rnorm <- function(n, seed) {
  with_seed(seed, {
    rnorm(n)
```

```

    })
}

pure_rnorm(n = 5, seed = 123)
pure_rnorm(n = 5, seed = 123) # Same result!

```

In Python, `numpy` provides a more modern, object-oriented way:

```

import numpy as np

# Create a random number generator instance with a
#   ↪ seed
rng = np.random.default_rng(seed=123)
print(rng.standard_normal(5))

# If we re-create the same generator, we get the
#   ↪ same numbers
rng2 = np.random.default_rng(seed=123)
print(rng2.standard_normal(5))

```

The key is the same: the “state” (the seed) is explicitly managed, not hidden globally.

6.2.2 The OOP Caveat in Python

This introduces a concept from OOP: the `rng` variable is an *object* that bundles together data (its internal seed state) and methods (`.standard_normal()`). This is **encapsulation**.

This is a double-edged sword for reproducibility. The `rng` object is now a stateful entity. If we called `rng.standard_normal(5)`

6 Functional Programming

a second time, it would produce different numbers because its internal state was mutated.

Core Python libraries like `pandas`, `scikit-learn`, and `matplotlib` are fundamentally object-oriented. Our guiding principle must be:

Use functions for the flow and logic of your analysis, and treat objects from libraries as values that are passed between these functions.

Avoid building your own complex classes with hidden state for your data pipeline. A pipeline composed of functions (`df2 = clean_data(df1); df3 = analyze_data(df2)`) is almost always more transparent than an OOP one (`pipeline.load(); pipeline.clean(); pipeline.analyze()`).

6.3 Writing Your Own Functions

Let's learn the syntax. The goal is always to encapsulate a single, logical piece of work.

6.3.1 In R

R functions are first-class citizens. You can assign them to variables and pass them to other functions.

```
# A simple function
calculate_ci <- function(x, level = 0.95) {
  se <- sd(x, na.rm = TRUE) / sqrt(length(x))
  mean_val <- mean(x, na.rm = TRUE)
```

```

alpha <- 1 - level
lower <- mean_val - qnorm(1 - alpha/2) * se
upper <- mean_val + qnorm(1 - alpha/2) * se

c(mean = mean_val, lower = lower, upper = upper)
}

# Use it
data <- c(1.2, 1.5, 1.8, 1.3, 1.6, 1.7)
calculate_ci(data)

```

For data analysis, you'll often want functions that work with column names. The `{dplyr}` package uses “tidy evaluation”:

```

library(dplyr)

summarize_variable <- function(dataset,
  ↪ var_to_summarize) {
  dataset %>%
    summarise(
      n = n(),
      mean = mean({{ var_to_summarize }}, na.rm =
        ↪ TRUE),
      sd = sd({{ var_to_summarize }}, na.rm = TRUE)
    )
}

# The {{ }} syntax tells dplyr to use the column
# name passed in
starwars %>%
  group_by(species) %>%
  summarize_variable(height)

```

6.3.2 In Python

Python uses the `def` keyword. Type hints are a best practice:

```
import pandas as pd

def summarize_variable_py(dataset: pd.DataFrame,
    ↵ var_to_summarize: str) -> pd.DataFrame:
    """Calculates summary statistics for a given
    ↵ column."""
    summary = dataset.groupby('species').agg(
        n=(var_to_summarize, 'size'),
        mean=(var_to_summarize, 'mean'),
        sd=(var_to_summarize, 'std'))
    ).reset_index()
    return summary
```

6.4 The Functional Toolkit: Map, Filter, and Reduce

Most `for` loops can be replaced by one of three core functional concepts: **mapping**, **filtering**, or **reducing**. These are “higher-order functions”—functions that take other functions as arguments.

6.4.1 1. Mapping: Applying a Function to Each Element

The pattern: You have a list of things, and you want to perform the same action on each element, producing a new list of

the same length.

6.4.1.1 In R with purrr::map()

The `{purrr}` package is the gold standard for functional programming in R:

- `map()`: Always returns a list
- `map_dbl()`: Returns a vector of doubles (numeric)
- `map_chr()`: Returns a vector of characters (strings)
- `map_lgl()`: Returns a vector of logicals (booleans)

```
library(purrr)

# The classic for-loop way (verbose)
means_loop <- vector("double", ncol(mtcars))
for (i in seq_along(mtcars)) {
  means_loop[[i]] <- mean(mtcars[[i]], na.rm = TRUE)
}

# The functional way with map_dbl()
means_functional <- map_dbl(mtcars, mean, na.rm =
  TRUE)
```

The `map()` version is not just shorter; it's safer. You can't make an off-by-one error.

6.4.1.2 In Python with List Comprehensions

Python's most idiomatic tool for mapping is the **list comprehension**:

6 Functional Programming

```
numbers = [1, 2, 3, 4, 5]
squares = [n**2 for n in numbers]
# > [1, 4, 9, 16, 25]
```

Python also has a built-in `map()` function:

```
def to_upper_case(s: str) -> str:
    return s.upper()

words = ["hello", "world"]
upper_words = list(map(to_upper_case, words))
# > ['HELLO', 'WORLD']
```

6.4.2 2. Filtering: Keeping Elements That Match a Condition

The pattern: You have a list of things, and you want to keep only the elements that satisfy a certain condition.

6.4.2.1 In R with `purrr::keep()`

```
df1 <- data.frame(x = 1:50)
df2 <- data.frame(x = 1:200)
df3 <- data.frame(x = 1:75)
list_of_dfs <- list(a = df1, b = df2, c = df3)

# Keep only data frames with more than 100 rows
large_dfs <- keep(list_of_dfs, ~ nrow(.x) > 100)
```

6.4.2.2 In Python with List Comprehensions

List comprehensions have a built-in if clause:

```
numbers = [1, 10, 5, 20, 15, 30]
large_numbers = [n for n in numbers if n > 10]
# > [20, 15, 30]
```

6.4.3 3. Reducing: Combining All Elements into a Single Value

The pattern: You have a list of things, and you want to iteratively combine them into a single summary value.

6.4.3.1 In R with purrr::reduce()

```
# Sum all elements
total_sum <- reduce(c(1, 2, 3, 4, 5), `+`)

# Find common columns across multiple data frames
list_of_colnames <- map(list_of_dfs, names)
common_cols <- reduce(list_of_colnames, intersect)
```

6.4.3.2 In Python with functools.reduce

6 Functional Programming

```
from functools import reduce
import operator

numbers = [1, 2, 3, 4, 5]
total_sum = reduce(operator.add, numbers)
# > 15
```

6.5 The Power of Composition

The final, beautiful consequence of a functional style is **composition**. You can chain functions together to build complex workflows from simple, reusable parts.

This R code is a sequence of function compositions:

```
starwars %>%
  filter(!is.na(mass)) %>%
  select(species, sex, mass) %>%
  group_by(sex, species) %>%
  summarise(mean_mass = mean(mass), .groups =
    ~ "drop")
```

The same idea in Python with pandas:

```
(starwars_py
  .dropna(subset=['mass'])
  .filter(items=['species', 'sex', 'mass'])
  .groupby(['sex', 'species'])
  ['mass'].mean()
  .reset_index()
)
```

Each step is a function that takes a data frame and returns a new, transformed data frame. By combining `map`, `filter`, and `reduce` with this compositional style, you can express complex data manipulation pipelines without writing a single `for` loop.

6.5.1 The Composition Challenge in Python

Method chaining in pandas is elegant but limited to the methods defined for `DataFrame` objects. R’s pipe operators (`|>` and `%>%`) are more flexible because functions are not strictly owned by objects—they can be more easily combined.

This reflects the languages’ different philosophies:

- R’s lineage traces back to Scheme (a Lisp dialect), making functional composition natural
- Python was designed around “one obvious way to do it,” favouring explicit R is fundamentally a functional language that acquired OOP features, while Python is an OOP language with powerful functional capabilities. Recognising this helps you leverage the native strengths of each.

6.6 Summary

This chapter has laid the groundwork for writing reproducible code by embracing **Functional Programming**.

Key takeaways:

- **Pure functions** guarantee the same output for the same input—no hidden dependencies on global state

6 Functional Programming

- Make impure operations (like randomness) explicit by controlling the seed
- Replace error-prone `for` loops with **map**, **filter**, and **reduce**
- Use **composition** to build complex pipelines from simple, reusable functions
- In Python, treat stateful library objects as values passed between pure functions

Understanding the distinction between R’s functional heritage and Python’s OOP nature is key to becoming an effective data scientist in either language. By mastering the functional paradigm, you’re building a foundation for code that is robust, easy to review, simple to debug, and truly reproducible.

In the next chapter, we’ll put these principles into practice with **{riexpress}**—a package that leverages functional composition and Nix to build fully reproducible, polyglot analytical pipelines.

7 Building Reproducible Pipelines with rixpress

7.1 Introduction: From Scripts and Notebooks to Pipelines

So far, we have learned about reproducible development environments with Nix and {rix}. We can now create project-specific environments with precise versions of R, Python, and all dependencies. But there's one more piece to the puzzle: **orchestration**.

How do we take our collection of functions and data files and run them in the correct order to produce our final data product? This problem of managing computational workflows is not new, and a whole category of **build automation tools** has been created to solve it.

7.1.1 The Evolution of Build Automation

The original solution, dating back to the 1970s, is **make**. Created by Stuart Feldman at Bell Labs in 1976, **make** reads a **Makefile** that describes the dependency graph of a project. If you change the code that generates `plot.png`, **make** is smart enough to only re-run the steps needed to rebuild the plot and the final report.

The strength of these tools is their language-agnosticism, but their weaknesses are twofold:

1. **File-centric**: You must manually handle all I/O—your first script saves `data.csv`, your second loads it. This adds boilerplate and surfaces for error.
2. **Environment-agnostic**: They track files but know nothing about the *software environment* needed to create those files.

This is where R’s `{targets}` package shines. It tracks dependencies between **R objects directly**, automatically handling serialisation. But `{targets}` operates within a single R session—for polyglot pipelines, you must manually coordinate via `{reticulate}`.

7.1.2 The Separation Problem

All these tools—from `make` to `{targets}` to Airflow—separate workflow management from environment management. You use one tool to run the pipeline and another (Docker, `{renv}`) to set up the software.

This separation creates friction. Running `{targets}` inside Docker ensures reproducibility, but forces the entire pipeline into one monolithic environment. What if your Python step requires TensorFlow 2.15 but your R step needs reticulate with Python 3.9? You’re stuck.

7.1.3 The Imperative Approach: Make + Docker

To illustrate this, consider the traditional setup for a polyglot pipeline. You’d need:

7.1 Introduction: From Scripts and Notebooks to Pipelines

1. A Dockerfile to set up the environment
2. A Makefile to orchestrate the workflow
3. Wrapper scripts for each step

Here's what a Makefile might look like:

```
# Makefile for a Python → R pipeline

DATA_DIR = data
OUTPUT_DIR = output

$(OUTPUT_DIR)/predictions.csv: $(DATA_DIR)/raw.csv
    ↳ scripts/train_model.py
        python scripts/train_model.py
    ↳ $(DATA_DIR)/raw.csv $@

$(OUTPUT_DIR)/plot.png:
    ↳ $(OUTPUT_DIR)/predictions.csv
    ↳ scripts/visualise.R
        Rscript scripts/visualise.R $< $@

$(OUTPUT_DIR)/report.html: $(OUTPUT_DIR)/plot.png
    ↳ report.qmd
        quarto render report.qmd -o $@

all: $(OUTPUT_DIR)/report.html

clean:
    rm -rf $(OUTPUT_DIR)/*
```

This looks clean, but notice the procedural boilerplate required in each script. Your `train_model.py` must parse command-line arguments and handle file I/O:

7 Building Reproducible Pipelines with rixpress

```
# scripts/train_model.py
import sys
import pandas as pd
from sklearn.ensemble import RandomForestClassifier

def main():
    input_path = sys.argv[1]
    output_path = sys.argv[2]

    # Load data
    df = pd.read_csv(input_path)

    # ... actual ML logic ...

    # Save results
    predictions.to_csv(output_path, index=False)

if __name__ == "__main__":
    main()
```

And your `visualise.R` script needs the same boilerplate:

```
# scripts/visualise.R
args <- commandArgs(trailingOnly = TRUE)
input_path <- args[1]
output_path <- args[2]

# Load data
predictions <- read.csv(input_path)

# ... actual visualisation logic ...

# Save plot
```

```
ggsave(output_path, plot)
```

The scientific logic is buried under file I/O scaffolding. And the environment? That's a separate 100+ line `Dockerfile` you must maintain.

7.1.4 rixpress: Unified Orchestration

`{rixpress}` solves this by using Nix not just as a package manager, but as the build automation engine itself. Each pipeline step is a **Nix derivation**—a hermetically sealed build unit.

Compare the same pipeline in `{rixpress}`:

```
library(rixpress)

list(
  rxp_py_file(
    name = raw_data,
    path = "data/raw.csv",
    read_function = "lambda x: pandas.read_csv(x)"
  ),

  rxp_py(
    name = predictions,
    expr = "train_model(raw_data)",
    user_functions = "functions.py"
  ),

  rxp_py2r(
    name = predictions_r,
    expr = predictions
)
```

7 Building Reproducible Pipelines with rixpress

```
) ,  
  
    rxp_r(  
        name = plot,  
        expr = visualise(predictions_r),  
        user_functions = "functions.R"  
    ),  
  
    rxp_qmd(  
        name = report,  
        qmd_file = "report.qmd"  
    )  
) |>  
    rxp_populate()
```

And your `functions.py` contains *only* the scientific logic:

```
# functions.py  
def train_model(df):  
    # ... pure ML logic, no file I/O ...  
    return predictions
```

The difference is stark:

Aspect	Make + Docker	rixpress
Files needed	Dockerfile, Makefile, wrapper scripts	gen-env.R, gen-pipeline.R, function files
I/O handling	Manual in every script	Automatic via encoders/decoders

Aspect	Make + Docker	rixpress
Dependencies	Explicit file rules	Inferred from object references
Environment	Separate Docker setup	Unified via <code>{rix}</code>
Expertise needed	Linux admin, Make syntax	R programming

This provides two key benefits:

1. **True Polyglot Pipelines:** Each step can have its own Nix environment. A Python step runs in a pure Python environment, an R step in an R environment, a Quarto step in yet another—all within the same pipeline.
2. **Deep Reproducibility:** Each step is cached based on the cryptographic hash of *all* its inputs: the code, the data, *and the environment*. Any change in dependencies triggers a rebuild. This is reproducibility at the build level, not just the environment level.

The interface is heavily inspired by `{targets}`, so you get the ergonomic, object-passing feel you’re used to, combined with the bit-for-bit reproducibility of the Nix build system.

7.2 What is rixpress?

`{rixpress}` streamlines creation of *micropipelines*—small-to-medium, single-machine analytic pipelines—by expressing a pipeline in idiomatic R while delegating build orchestration to the Nix build system.

Key features:

7 Building Reproducible Pipelines with rixpress

- Define pipeline *derivations* with concise `rxp_*`() helper functions
- Seamlessly mix R, Python, Julia, and Quarto steps
- Reuse hermetic environments defined via `{rix}` and a `default.nix`
- Visualise and inspect the DAG; selectively read, load, or copy outputs
- Automatic caching: only rebuild what changed

Here is what a basic pipeline looks like:

```
library(rixpress)

list(
  rxp_r_file(
    mtcars,
    'mtcars.csv',
    \((x) read.csv(file = x, sep = "|"))
  ),

  rxp_r(
    mtcars_am,
    filter(mtcars, am == 1)
  ),

  rxp_r(
    mtcars_head,
    head(mtcars_am)
  ),

  rxp_qmd(
    page,
    "page.qmd"
)
```

```
) |>  
  rxp_populate()
```

7.3 Getting Started

7.3.1 Initialising a project

If you're starting fresh, you can bootstrap a project using a temporary shell:

```
nix-shell --expr "$(curl -sl  
  https://raw.githubusercontent.com/ropensci/rix/main/inst,
```

Once inside, start R and run:

```
rixpress::rxp_init()
```

This creates two essential files:

- **gen-env.R**: Where you define your environment with {rix}
- **gen-pipeline.R**: Where you define your pipeline with {rixpress}

7.3.2 Defining the environment

Open `gen-env.R` and define the tools your pipeline needs:

7 Building Reproducible Pipelines with rixpress

```
library(rix)

rix(
  date = "2025-10-14",
  r_pkgs = c("dplyr", "ggplot2", "quarto",
    ↪ "rixpress"),
  ide = "none",
  project_path = ".",
  overwrite = TRUE
)
```

Run this script to generate `default.nix`, then build and enter your environment:

```
nix-build  
nix-shell
```

7.3.3 Defining the pipeline

Open `gen-pipeline.R` and define your pipeline:

```
library(rixpress)

list(
  rxp_r_file(
    name = mtcars,
    path = "data/mtcars.csv",
    read_function = \((x) read.csv(x, sep = "|"))
  ),
  rxp_r(
```

```

  name = mtcars_am,
  expr = dplyr::filter(mtcars, am == 1)
),

rxp_r(
  name = mtcars_head,
  expr = head(mtcars_am)
)
) |>
  rxp_populate()

```

Running `rxp_populate()` generates a `pipeline.nix` file and builds the entire pipeline.

7.4 Core Functions

7.4.1 Defining derivations

`{rixpress}` provides several functions to define pipeline steps:

Function	Purpose
<code>rxp_r()</code>	Run R code
<code>rxp_r_file()</code>	Read a file using R
<code>rxp_py()</code>	Run Python code
<code>rxp_py_file()</code>	Read a file using Python
<code>rxp_qmd()</code>	Render a Quarto document
<code>rxp_py2r()</code>	Convert Python object to R
<code>rxp_r2py()</code>	Convert R object to Python

7.4.2 Building the pipeline

```
# Generate pipeline.nix only (don't build)
rxp_populate(build = FALSE)

# Build the pipeline
rxp_make()
```

7.4.3 Inspecting outputs

Because outputs live in `/nix/store/`, `{rixpress}` provides helpers:

```
# List all built artifacts
rxp_inspect()

# Read an artifact into R
result <- rxp_read("mtcars_head")

# Load an artifact into the global environment
rxp_load("mtcars_head")

# Copy an output file to current directory
rxp_copy("page")
```

7.4.4 Visualizing the pipeline

```
# Static DAG plot
rxp_ggdag()

# Interactive network
rxp_visnetwork()

# Text-based trace
rxp_trace()
```

7.5 Polyglot Pipelines

One of {rixpress}'s strengths is seamlessly mixing languages. Here's a pipeline that reads data with Python's `polars`, processes it with R's `dplyr`, and renders a Quarto report:

```
library(rixpress)

list(
  rxp_py_file(
    name = mtcars_pl,
    path = "data/mtcars.csv",
    read_function = "lambda x: polars.read_csv(x,
      ↵  separator='|')",
  ),
  rxp_py(
    name = mtcars_filtered,
    expr = "mtcars_pl.filter(polars.col('am') ==
      ↵  1).to_pandas()",
  ),
)
```

```
rxp_py2r(  
  name = mtcars_r,  
  expr = mtcars_filtered  
)  
  
rxp_r(  
  name = mtcars_head,  
  expr = head(mtcars_r)  
)  
  
rxp_qmd(  
  name = report,  
  qmd_file = "report.qmd"  
)  
) |>  
  rxp_populate()
```

7.5.1 Method 1: Using language converters

The `rxp_py2r()` and `rxp_r2py()` functions use `{reticulate}` to convert objects between languages:

```
rxp_py2r(  
  name = mtcars_r,  
  expr = mtcars_py  
)
```

7.5.2 Method 2: Using universal data formats

For more control, use `encoder` and `decoder` arguments to serialize to formats like JSON:

```
# Python step: serialize to JSON
rxp_py(
    name = mtcars_json,
    expr = "mtcars_pl.filter(polars.col('am') == 1)",
    user_functions = "functions.py",
    encoder = "serialize_to_json"
),

# R step: deserialize from JSON
rxp_r(
    name = mtcars_head,
    expr = my_head(mtcars_json),
    user_functions = "functions.R",
    decoder = "jsonlite::fromJSON"
)
```

This approach makes your pipeline more modular—any language that can read JSON could be added in the future.

7.6 Caching and Incremental Builds

One of the most powerful features of using Nix for pipelines is automatic caching. Because Nix tracks all inputs to each derivation, it knows exactly what needs to be rebuilt when something changes.

Try this:

1. Build your pipeline with `rxp_make()`
2. Change one step in your pipeline
3. Run `rxp_make()` again

Nix will detect that unchanged steps are already cached and instantly reuse them. It only rebuilds the steps affected by your change.

7.7 Build Logs and Debugging

Every time you run `rxp_populate()`, a timestamped log is saved in the `_rixpress/` directory. This is like having a Git history for your pipeline's *outputs*.

```
# List all past builds
rxp_list_logs()

# Load artifact from current build
new_result <- rxp_read("mtcars_head")

# Load artifact from previous build
old_result <- rxp_read("mtcars_head", which_log =
  ↵ "z9y8x")

# Compare them
identical(new_result, old_result)
```

This is incredibly powerful for debugging and validation—you can go back in time to inspect any output from any previous pipeline run.

7.8 `ryxpress`: The Python Interface

If you prefer working in Python, `ryxpress` provides the same functionality. You still define your pipeline in R (since that's where `{rixpress}` runs), but you can build and inspect artifacts from Python.

To set up an environment with `ryxpress`:

```
rix(
  date = "2025-10-14",
  r_pkgs = c("rixpress"),
  py_conf = list(
    py_version = "3.13",
    py_pkgs = c("ryxpress", "rds2py", "biocframe",
               "pandas")
  ),
  ide = "none",
  project_path = ".",
  overwrite = TRUE
)
```

Then from Python:

```
from ryxpress import rxp_make, rxp_inspect, rxp_load

# Build the pipeline
rxp_make()

# Inspect artifacts
rxp_inspect()

# Load an artifact
```

```
r xp_load("mtcars_head")
```

rixpress handles the conversion automatically:

- Tries `pickle.load` first
- Falls back to `rds2py` for R objects
- Returns file paths for complex outputs

7.9 Running Someone Else's Pipeline

The ultimate test of reproducibility: can someone else run your pipeline?

With a Nix-based workflow, they need only:

1. `git clone` your repository
2. Run `nix-build && nix-shell`
3. Run `source("gen-pipeline.R")` or `r xp_make()`

That's it. Nix reads your `default.nix` and `pipeline.nix` files and builds the *exact* same environment and data product, bit-for-bit.

7.10 Exporting and Importing Artifacts

For CI/CD or sharing between machines:

```
# Export build products to a tarball
rxp_export_artifacts()

# Import on another machine before building
rxp_import_artifacts()
```

This speeds up continuous integration by avoiding unnecessary rebuilds.

7.11 Real-World Examples

The rixpress_demos repository contains many complete examples. Here are a few patterns worth studying.

7.11.1 Example 1: Machine Learning with XGBoost

This pipeline trains an XGBoost classifier in Python, then passes predictions to R for evaluation with {yardstick}:

```
library(rixpress)

list(
  # Load data as NumPy array
  rxp_py_file(
    name = dataset_np,
    path = "data/pima-indians-diabetes.csv",
    read_function = "lambda x: loadtxt(x,
      delimiter=',')"
```

7 Building Reproducible Pipelines with rixpress

```
),  
  
# Split features and target  
rxp_py(name = X, expr = "dataset_np[:,0:8]"),  
rxp_py(name = Y, expr = "dataset_np[:,8]"),  
  
# Train/test split  
rxp_py(  
    name = splits,  
    expr = "train_test_split(X, Y, test_size=0.33,  
        random_state=7)"  
)  
  
# Extract splits  
  
rxp_py(name = X_train, expr = "splits[0]"),  
rxp_py(name = X_test, expr = "splits[1]"),  
rxp_py(name = y_train, expr = "splits[2]"),  
rxp_py(name = y_test, expr = "splits[3]"),  
  
# Train XGBoost model  
rxp_py(  
    name = model,  
    expr = "XGBClassifier(use_label_encoder=False,  
        eval_metric='logloss').fit(X_train,  
        y_train)"  
)  
  
# Make predictions  
rxp_py(name = y_pred, expr =  
    "model.predict(X_test)"),  
  
# Export predictions to CSV for R
```

```
rxp_py(  
    name = combined_df,  
    expr = "DataFrame({'truth': y_test, 'estimate':  
        ~ y_pred})")  
,  
  
rxp_py(  
    name = combined_csv,  
    expr = "combined_df",  
    user_functions = "functions.py",  
    encoder = "write_to_csv"  
,  
  
# Compute confusion matrix in R  
rxp_r(  
    combined_factor,  
    expr = mutate(combined_csv, across(everything(),  
        ~ factor)),  
    decoder = "read.csv"  
,  
  
rxp_r(  
    name = confusion_matrix,  
    expr = yardstick::conf_mat(combined_factor,  
        ~ truth, estimate)  
)  
) |>  
    rxp_populate(build = FALSE)  
  
# Adjust Python imports  
adjust_import("import numpy", "from numpy import  
    ~ array, loadtxt")  
adjust_import("import xgboost", "from xgboost import  
    ~ XGBClassifier")
```

7 Building Reproducible Pipelines with rixpress

```
adjust_import("import sklearn", "from
←   sklearn.model_selection import
←   train_test_split")
add_import("from pandas import DataFrame",
←   "default.nix")

rxp_make()
```

This demonstrates:

- Python-heavy computation with XGBoost
- Custom serialization via `encoder/decoder`
- Adjusting Python imports with `adjust_import()` and `add_import()`
- Passing results to R for evaluation

7.11.2 Example 2: Reading Many Input Files

When you have multiple CSV files in a directory:

```
library(rixpress)

list(
  # R approach: read all files at once
  rxp_r_file(
    name = mtcars_r,
    path = "data",
    read_function = \ (x) {
      readr::read_delim(list.files(x, full.names =
        TRUE), delim = "|")
    }
  ),
)
```

```
# Python approach: custom function
rxp_py_file(
  name = mtcars_py,
  path = "data",
  read_function = "read_many_csvs",
  user_functions = "functions.py"
),

rxp_py(
  name = head_mtcars,
  expr = "mtcars_py.head()"
)
) |>
  rxp_populate()
```

The key insight: `r xp_r_file()` and `r xp_py_file()` can point to a directory, and your `read_function` handles the logic.

7.11.3 Example 3: Full Python→R→Quarto Workflow

A complete pipeline that bounces data between languages and renders a report:

```
library(rixpress)

list(
  # Read with Python polars
  rxp_py_file(
    name = mtcars_pl,
```

7 Building Reproducible Pipelines with rixpress

```
path = "data/mtcars.csv",
read_function = "lambda x: polars.read_csv(x,
    ↵ separator='|')"

),

# Filter in Python, convert to pandas for
#   ↵ reticulate
rxp_py(
    name = mtcars_pl_am,
    expr = "mtcars_pl.filter(polars.col('am') ==
        ↵ 1).to_pandas()"
),

# Convert to R
rxp_py2r(name = mtcars_am, expr = mtcars_pl_am),

# Process in R
rxp_r(
    name = mtcars_head,
    expr = my_head(mtcars_am),
    user_functions = "functions.R"
),

# Back to Python
rxp_r2py(name = mtcars_head_py, expr =
    ↵ mtcars_head),

# More Python processing
rxp_py(name = mtcars_tail_py, expr =
    ↵ "mtcars_head_py.tail()"),

# Back to R
rxp_py2r(name = mtcars_tail, expr =
    ↵ mtcars_tail_py),
```

```
# Final R step
rxp_r(name = mtcars_mpg, expr =
  ↴ dplyr::select(mtcars_tail, mpg)),

# Render Quarto document
rxp_qmd(
  name = page,
  qmd_file = "my_doc/page.qmd",
  additional_files = c("my_doc/content.qmd",
    ↴ "my_doc/images")
)
) |>
  rxp_populate()
```

Note the `additional_files` argument for `r xp_qmd()`: this includes child documents and images that the main Quarto file needs.

7.11.4 More Examples

The rixpress_demos repository includes:

- **jl_example**: Using Julia in pipelines
- **r_qs**: Using `{qs}` for faster serialization
- **python_r_typst**: Compiling to Typst documents
- **r_multi_envs**: Different Nix environments for different derivations
- **yanai_lercher_2020**: Reproducing a published paper's analysis

7.12 Summary

`{riexpress}` unifies environment management and workflow orchestration:

- **Define** pipelines with `rxp_*`() functions in familiar R syntax
- **Mix** languages freely—R, Python, Julia, Quarto
- **Build** with Nix for deterministic, cached execution
- **Inspect** outputs with `rxp_read()`, `rxp_load()`, `rxp_copy()`
- **Debug** with timestamped build logs
- **Share** reproducible pipelines via Git

The Python port `ryxpress` provides the same experience for Python-first workflows.

By embracing structured, plain-text pipelines over notebooks for production work, your analysis becomes more reliable, more scalable, and fundamentally more reproducible.

8 Containerisation With Docker

8.1 Introduction

Up until now, we've been using Nix as a powerful tool for creating reproducible development environments directly on our machines. Nix gives us fine-grained control over every package and dependency in our project, ensuring bit-for-bit reproducibility. However, when it comes to distributing a *data product*, another technology, Docker, is incredibly popular.

While Nix manages dependencies for an application that runs on a host operating system, Docker takes a different approach: it packages an application *along with* a lightweight operating system and all its dependencies into a single, portable unit called a **container**. This container can then run on any machine that has Docker installed, regardless of its underlying OS.

8.1.1 Spatial vs Temporal Reproducibility

To understand when and why to use Docker, it helps to distinguish between two types of reproducibility:

- **Spatial reproducibility:** The ability to execute an analysis identically across different machines *right now*. Docker excels here—a container runs the same way on your laptop, a colleague’s workstation, and a cloud server.
- **Temporal reproducibility:** The ability to execute an analysis identically *over time*. Docker is weaker here. The imperative commands in a `Dockerfile` (like `apt-get update`) are non-deterministic—rebuilding the same file at different times can yield different images.

This distinction is crucial. Docker’s true reproducibility promise is that a specific, pre-built image will always launch an identical container. It does *not* promise that building the same `Dockerfile` twice will yield an identical image.

As empirical research has shown, achieving deterministic builds requires systems designed for this purpose—like Nix’s functional model, where identical inputs always produce identical outputs. Studies rebuilding historical Nix packages found bit-for-bit reproducibility rates exceeding 90%.

8.1.2 The Best of Both Worlds

This is why we combine Nix and Docker rather than choosing one or the other:

- **Nix** provides strong temporal reproducibility—the guarantee that an environment built today will be identical when rebuilt years later.
- **Docker** provides strong spatial reproducibility and universal distribution—the guarantee that a container runs the same everywhere Docker runs.

The pattern we'll use is simple: use Nix *inside* a Docker container. Start with a minimal base image that has Nix installed. Then, use Nix to declaratively build the precise environment within the image. Docker becomes a portable runtime for this Nix-managed environment, excellent for deployment to systems where Nix isn't installed.

8.1.3 Interactive Development vs Distribution

This approach also clarifies *when* to use each tool:

- **Use Nix directly** for interactive development. It integrates seamlessly with your IDE and filesystem. No volume mounts, no port forwarding, no graphical application headaches.
- **Use Docker** for distributing finished data products. Package your `{rixpress}` pipeline into an image that anyone can run with a single command.

If you've never heard of Docker before, this chapter will provide the basic knowledge required to get started. Let's start by watching this very short video that introduces the core concepts.

The Rocker Project provides a large collection of ready-to-use Docker images for R users.

8.2 Docker Essentials

8.2.1 Installing Docker

The first step is to install Docker. You'll find the instructions for Ubuntu here, for Windows here (read the system requirements

8 Containerisation With Docker

section as well!) and for macOS here (make sure to choose the right version for the architecture of your Mac, if you have an M1 Mac use *Mac with Apple silicon*).

After installation, it might be a good idea to restart your computer, if the installation wizard does not invite you to do so. To check whether Docker was installed successfully, run the following command in a terminal:

```
docker run --rm hello-world
```

This should print the following message:

```
Hello from Docker!  
This message shows that your installation appears to  
be working correctly.
```

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
(amd64)
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

```
$ docker run -it ubuntu bash
```

Share images, automate workflows, and more with a free Docker ID:

<https://hub.docker.com/>

For more examples and ideas, visit:

<https://docs.docker.com/get-started/>

If you see this message, congratulations, you are ready to run Docker. If you see an error message about permissions, this means that something went wrong. If you're running Linux, make sure that your user is in the Docker group by running:

```
groups $USER
```

You should see your username and a list of groups that your user belongs to. If a group called `docker` is not listed, then you should add yourself to the group by following these steps.

8.2.2 The Rocker Project and Image Registries

When running a command like:

```
docker run --rm hello-world
```

what happens is that an image, in this case `hello-world`, gets pulled from a so-called *registry*. A registry is a storage and distribution system for Docker images. Think of it as a GitHub for Docker images, where you can push and pull images, much like you would with code repositories. The default public registry that Docker uses is called Docker Hub, but companies can also host their own private registries to store proprietary images.

8 Containerisation With Docker

Many open source projects build and distribute Docker images through Docker Hub, for example the Rocker Project.

The Rocker Project is instrumental for R users that want to use Docker. The project provides a large list of images that are ready to run with a single command. As an illustration, open a terminal and paste the following line:

```
docker run --rm -e PASSWORD=yourpassword -p
← 8787:8787 rocker/rstudio
```

Once this stops running, go to `http://localhost:8787/` and enter `rstudio` as the username and `yourpassword` as the password. You should login to a RStudio instance: this is the web interface of RStudio that allows you to work with R from a server. In this case, the *server* is the Docker container running the image. Yes, you've just pulled a Docker image containing Ubuntu with a fully working installation of RStudio web!

Let's open a new script and run the following lines:

```
data(mtcars)
summary(mtcars)
```

You can now stop the container (by pressing **CTRL-C** in the terminal). Let's now rerun the container... you should realise that your script is gone! This is the first lesson: whatever you do inside a container will disappear once the container is stopped. This also means that if you install the R packages that you need while the container is running, you will need to reinstall them every time.

Thankfully, the Rocker Project provides a list of images with many packages already available. For example to run R with the `{tidyverse}` collection of packages already pre-installed, run:

```
docker run --rm -e PASSWORD=yourpassword -p  
↪ 8787:8787 rocker/tidyverse
```

8.2.3 Basic Docker Workflow

You already know about running containers using `docker run`. With the commands we ran before, your terminal will need to stay open, or else, the container will stop. Starting now, we will run Docker commands in the background using the `-d` flag (*d* as in *detach*):

```
docker run --rm -d -e PASSWORD=yourpassword -p  
↪ 8787:8787 rocker/tidyverse
```

You can run several containers in the background simultaneously. List running containers with `docker ps`:

```
docker ps  
CONTAINER ID   IMAGE          COMMAND      CREATED      NAMES  
↪ STATUS       PORTS          NAMES  
c956fbbeebcb  rocker/tidyverse  "/init"    3 minutes ago  Up 3 minutes  
↪ 0.0.0.0:8787->8787/tcp    elastic_morse
```

Stop the container using its ID:

```
docker stop c956fbbeebcb
```

Let's discuss the other flags:

8 Containerisation With Docker

- **--rm**: Removes the container once it's stopped
- **-e**: Provides environment variables to the container (e.g., **PASSWORD**)
- **-p**: Sets the port mapping (host:container)
- **--name**: Gives the container a custom name

Run a container with a name:

```
docker run -d --name my_r --rm -e  
    ↳ PASSWORD=yourpassword -p 8787:8787  
    ↳ rocker/tidyverse
```

You can now interact with this container using its name:

```
docker exec -ti my_r bash
```

You are now inside a terminal session, inside the running container! This can be useful for debugging purposes.

Finally, let's solve the issue of our scripts disappearing. Create a folder somewhere on your computer, then run:

```
docker run -d --name my_r --rm -e  
    ↳ PASSWORD=yourpassword -p 8787:8787 \  
    ↳ -v  
        ↳ /path/to/your/local/folder:/home/rstudio/scripts:rw  
    ↳ rocker/tidyverse
```

You should now be able to save scripts inside the **scripts/** folder from RStudio and they will appear in the folder you created on your host machine.

8.3 Making Our Own Images

To create our own images, you can start from an image provided by an open source project like Rocker, or you can start from the base Ubuntu image. Since we are using Nix to set up the reproducible development environment, we can use `ubuntu:latest`—our development environment will always be exactly the same, thanks to Nix.

8.3.1 A Minimal Dockerfile with Nix

```
FROM ubuntu:latest

RUN apt update -y
RUN apt install curl -y

# Download the default.nix that comes with {rix}
RUN curl -O
↳ https://raw.githubusercontent.com/ropensci/rix/main/inst...

# Install Nix via Determinate Systems installer
RUN curl --proto '=https' --tlsv1.2 -sSf -L
↳ https://install.determinate.systems/nix | sh -s
↳ -- install linux \
--extra-conf "sandbox = false" \
--init none \
--no-confirm

# Add Nix to the path
ENV PATH="${PATH}:/nix/var/nix/profiles/default/bin"
ENV user=root
```

8 Containerisation With Docker

```
# Configure rstats-on-nix cache
RUN mkdir -p /root/.config/nix && \
    echo "substituters = https://cache.nixos.org
        https://rstats-on-nix.cachix.org" >
    /root/.config/nix/nix.conf && \
    echo "trusted-public-keys =
        cache.nixos.org-1:6NCHdD59X431o0gWypbMrAURkbJ16ZPMQFGsp
        rstats-on-nix.cachix.org-1:vdiiVgocg6WeJr0DIqdprZRUrhi" >> /root/.config/nix/nix.conf

# Copy script to generate environment
COPY gen-env.R .

# Generate default.nix from gen-env.R
RUN nix-shell --run "Rscript gen-env.R"

# Build the environment
RUN nix-build

# Run nix-shell when container starts
CMD nix-shell
```

Every **Dockerfile** starts with a **FROM** statement specifying the base image. Then, every command starts with **RUN**. We install and configure Nix, copy an R script to generate the environment, and then build it. Finally, we run **nix-shell** when executing a container (the **CMD** statement).

8.3.2 Splitting Into a Reusable Base Image

Because the Nix installation step is generic, we can split this into two stages. First, create a base image with just Nix installed:

```
# nix-base/Dockerfile
FROM ubuntu:latest AS nix-base

RUN apt update -y && apt install -y curl

# Install Nix via Determinate Systems installer
RUN curl --proto '=https' --tlsv1.2 -sSf -L
    https://install.determinate.systems/nix | sh -s
    -- install linux \
    --extra-conf "sandbox = false" \
    --init none \
    --no-confirm

ENV PATH="/nix/var/nix/profiles/default/bin:${PATH}"
ENV user=root

# Configure Nix binary cache
RUN mkdir -p /root/.config/nix && \
    echo "substituters = https://cache.nixos.org
        https://rstats-on-nix.cachix.org" >
    /root/.config/nix/nix.conf && \
    echo "trusted-public-keys =
        cache.nixos.org-1:6NCHdD59X431o0gWypbMrAURkbJ16ZPMQF0
        rstats-on-nix.cachix.org-1:vdiiVgocg6WeJrODIqdprZRUrh
        >> /root/.config/nix/nix.conf
```

Build and tag this image:

8 Containerisation With Docker

```
docker build -t nix-base:latest .
```

Now, for any project, simply reuse it:

```
FROM nix-base:latest

COPY gen-env.R .

RUN curl -O
  https://raw.githubusercontent.com/ropensci/rix/main/install.R
RUN nix-shell --run "Rscript gen-env.R"
RUN nix-build

CMD ["nix-shell"]
```

With `gen-env.R`:

```
library(rix)

rix(
  date = "2025-08-04",
  r_pkgs = c("dplyr", "ggplot2"),
  py_conf = list(
    py_version = "3.13",
    py_pkgs = c("polars", "great-tables")
  ),
  ide = "none",
  project_path = ".",
  overwrite = TRUE
)
```

Build and run:

```
docker build -t my-project .
docker run -it --rm --name my-project-container
    ↵ my-project
```

This drops you in an interactive Nix shell running inside Docker!

8.4 Publishing Images on Docker Hub

If you want to share Docker images through Docker Hub, you first need to create a free account. A free account gives you unlimited public repositories.

List all images on your computer:

```
docker images
```

Log in to Docker Hub:

```
docker login
```

Tag the image with your username:

```
docker tag IMAGE_ID your_username/nix-base:latest
```

Push the image:

8 Containerisation With Docker

```
docker push your_username/nix-base:latest
```

This image can now be used as a stable base for other projects:

```
FROM your_username/nix-base:latest
```

```
RUN mkdir ...
```

8.4.1 Sharing Without Docker Hub

If you can't upload to Docker Hub, you can save the image to a file:

```
docker save nix-base | gzip > nix-base.tgz
```

Load it on another machine:

```
gzip -d nix-base.tgz  
docker load < nix-base.tar
```

8.5 What If You Don't Use Nix?

Using Nix inside of Docker makes it very easy to set up an environment, but what if you can't use Nix for some reason? In this case, you would need to use other tools to install the right R or Python packages and it is likely going to be more difficult. The main issue you will face is missing development libraries.

For example, to install and use the R `{stringr}` package, you will need to first install `libicu-dev`:

```
FROM rocker/r-ver:4.5.1

RUN apt-get update && apt-get install -y \
    libglpk-dev \
    libxml2-dev \
    libcairo2-dev \
    libgit2-dev \
    libcurl4-openssl-dev \
    # ... many more
```

Another issue is that building the image is not a reproducible process, only running containers is. To mitigate this, use tagged images or better yet, a digest:

```
FROM
↳ rocker/r-ver@sha256:1dbe7a6718b7bd8630addc45a32731624fb7b
```

This will always pull exactly the same layers. However, at some point, that version of Ubuntu will be outdated. Using Nix, you can stay on `ubuntu:latest`.

8.6 Building Docker Images Directly with Nix

So far, we've used Nix *inside* Docker containers. But there's an even more powerful approach: using Nix to build Docker images directly, bypassing `Dockerfile`s entirely.

Nix provides `dockerTools`, a set of functions that can create OCI-compliant container images. Because these images are

8 Containerisation With Docker

built through Nix’s deterministic build system, they have stronger reproducibility guarantees than images built with `docker build`.

8.6.1 Why Skip the Dockerfile?

A `Dockerfile` is fundamentally imperative—it’s a script of commands executed in order. Commands like `apt-get update` are non-deterministic by nature. Even with careful pinning, you’re fighting the tool’s design.

Nix’s `dockerTools.buildImage` is declarative. You describe what should be in the image, and Nix figures out how to build it reproducibly. The resulting image is a pure function of its inputs.

8.6.2 A Basic Example

Here’s a Nix expression that builds a Docker image containing R and some packages:

```
# docker-image.nix
let
  pkgs = import (fetchTarball
    ↗ "https://github.com/rstats-on-nix/nixpkgs/archive/2025-
  ) {};

  # Define the R environment
  myR = pkgs.rWrapper.override {
    packages = with pkgs.rPackages; [
      dplyr
```

```

ggplot2
quarto
];
};

in pkgs.dockerTools.buildImage {
  name = "my-r-env";
  tag = "latest";

  copyToRoot = pkgs.buildEnv {
    name = "image-root";
    paths = [ myR pkgs.quarto pkgs.coreutils
              ↳ pkgs.bash ];
    pathsToLink = [ "/bin" ];
  };

  config = {
    Cmd = [ "${pkgs.bash}/bin/bash" ];
    WorkingDir = "/work";
  };
}

```

Build and load it:

```

# Build the image (outputs a .tar.gz file)
nix-build docker-image.nix

# Load into Docker
docker load < result

# Run it
docker run -it --rm my-r-env:latest

```

8.6.3 Benefits Over Dockerfiles

Aspect	Dockerfile	Nix dockerTools
Reproducibility	Hyperative, non-deterministic	Declarative, deterministic
Layer caching	Based on command order	Based on content hashes
Image size	Often includes unnecessary packages	Minimal—only what you specify
Composition	Copy-paste between files	Reuse Nix expressions

8.6.4 Layered Images for Efficiency

For faster CI builds, you can create layered images where the base environment is cached:

```
pkgs.dockerTools.buildLayeredImage {
    name = "my-analysis";
    tag = "latest";

    contents = [ myR pkgs.quarto ];

    config = {
        Cmd = [ "${myR}/bin/R" "--vanilla" ];
    };
}
```

`buildLayeredImage` creates separate layers for each package, so unchanged dependencies are cached between builds.

8.6.5 When to Use This Approach

- **CI/CD pipelines:** When you need reproducible image builds as part of automated workflows
- **Minimal images:** When image size matters and you want only what you need
- **Complex environments:** When the environment is already defined in Nix and you want to deploy it as a container

For simpler use cases, the “Nix inside Docker” approach from earlier sections may be more accessible.

8.7 Dockerizing a rixpress Pipeline

We can package our entire `{rixpress}` project into a single Docker image. This image can then be run by anyone with Docker installed, regardless of their host operating system or whether they have Nix.

Assume your project directory has:

```
.  
  data/  
    mtcars.csv  
  gen-env.R  
  gen-pipeline.R  
  functions.R  
  functions.py  
  default.nix  
  pipeline.nix
```

8.7.1 Step 1: The Dockerfile

```
FROM nix-base:latest
# Or: FROM your-username/nix-base:latest

WORKDIR /app

# Copy all project files
COPY . .

# Build the pipeline during image build
RUN nix-build pipeline.nix

# Export results when container runs
COPY export-results.R .
CMD ["nix-shell", "--run", "Rscript
    ↪ export-results.R"]
```

8.7.2 Step 2: The Export Script

```
# export-results.R
library(rixpress)
library(jsonlite)

output_dir <- "/output"
dir.create(output_dir, showWarnings = FALSE)

message("Reading target 'mtcars_head'...")
final_data <- rxp_read("mtcars_head")
```

```
output_path <- file.path(output_dir,  
    ↵ "mtcars_analysis_result.json")  
write_json(final_data, output_path, pretty = TRUE)  
  
message(paste("Successfully exported result to",  
    ↵ output_path))
```

8.7.3 Step 3: Build and Run

Build:

```
docker build -t my-reproducible-pipeline .
```

Run to get results:

```
mkdir -p ./output  
  
docker run --rm --name my_pipeline_run \  
    -v "$(pwd)/output":/output \  
    my-reproducible-pipeline
```

Check your local `output` directory—you’ll find the `mtcars_analysis_res` file containing the exact, reproducible result of your pipeline.

You have successfully packaged a complex, polyglot pipeline into a simple, portable Docker image. This workflow combines the best of both worlds: Nix’s power for creating reproducible builds and Docker’s universal standard for distributing and running applications.

8.8 Summary

Docker and Nix complement each other:

- **Docker** provides a universal runtime and distribution mechanism
- **Nix** provides bit-for-bit reproducible environment management
- **Together** they enable truly reproducible, portable data products

Key Docker concepts:

- **Images** are templates; **containers** are running instances
- Containers are ephemeral—use **volumes** for persistence
- Use **registries** (Docker Hub) to share images
- The **Rocker Project** provides ready-made R images

For reproducibility:

- Use `ubuntu:latest` + Nix rather than pinning specific OS versions
- Build a reusable `nix-base` image to share across projects
- Package `{riexpress}` pipelines for distribution

8.9 Further Reading

- Running Your R Script in Docker
- Docker & R Reproducibility
- R Docker Tutorial

9 Robust Pipelines with Monads

In the previous chapters, you learned functional programming fundamentals and how to build reproducible pipelines with `{rixpress}`. Now we'll add another layer of robustness: **monads**.

Monads might sound abstract, but they solve concrete problems involving **Logging** (tracing what each step does without cluttering your functions), **Error handling** (letting failures propagate gracefully instead of crashing), and **Composition** (keeping functions composable even when they need to do “extra” work).

By the end of this chapter, you'll know how to integrate `{chronicler}` and `talvez` into your pipelines for more robust, observable data workflows.

9.1 The Problem: Decorated Functions Don't Compose

Suppose you want your functions to provide logs. You might rewrite `sqrt()` like this:

9 Robust Pipelines with Monads

```
my_sqrt <- function(x, log = "") {  
  list(  
    result = sqrt(x),  
    log = c(log, paste0("Running sqrt with input ",  
      ↵  x)))  
}  
  
my_log <- function(x, log = "") {  
  list(  
    result = log(x),  
    log = c(log, paste0("Running log with input ",  
      ↵  x)))  
}
```

These functions now return lists with both the result and a log. But there's a problem—they don't compose:

```
# This works:  
10 |> sqrt() |> log()  
  
# This fails:  
10 |> my_sqrt() |> my_log()  
# Error: non-numeric argument to mathematical  
  ↵  function
```

`my_log()` expects a number, but `my_sqrt()` returns a list. We've broken composition.

9.2 The Solution: Function Factories and Bind

A monad provides two things:

1. **A function factory** that decorates functions so they can provide additional output without rewriting their core logic
2. **A bind() function** that makes these decorated functions compose

Here's a simple function factory for logging:

```
log_it <- function(.f) {  
  fstring <- deparse(substitute(.f))  
  
  function(..., .log = NULL) {  
    list(  
      result = .f(...),  
      log = c(.log, paste0("Running ", fstring, "  
        with argument ", ...))  
    )  
  }  
}  
  
# Create decorated functions  
l_sqrt <- log_it(sqrt)  
l_log <- log_it(log)  
  
l_sqrt(10)  
#> $result  
#> [1] 3.162278  
#> $log  
#> [1] "Running sqrt with argument 10"
```

Now we need `bind()` to make them compose:

```
bind <- function(.l, .f) {  
  .f(.l$result, .log = .l$log)  
}  
  
# Now they compose!  
10 |>  
  l_sqrt() |>  
  bind(l_log)  
  
#> $result  
#> [1] 1.151293  
#> $log  
#> [1] "Running sqrt with argument 10"  
#> [2] "Running log with argument 3.16227766016838"
```

This pattern of a function factory plus `bind()` is the essence of a monad.

9.3 The `chronicler` Package

The `{chronicler}` package implements this pattern properly for R. It provides:

- `record()`: A function factory that decorates functions
- `bind_record()`: The bind operation
- Automatic logging of all operations

```
library(chronicler)

# Decorate functions
r_sqrt <- record(sqrt)
r_exp <- record(exp)
r_mean <- record(mean)

# Compose them
result <- 1:10 |>
  r_sqrt() |>
  bind_record(r_exp) |>
  bind_record(r_mean)

# View the result
result$value
#> [1] 5.187899

# View the log
read_log(result)
#> [1] "Complete log:"
#> [2] " sqrt ran successfully"
#> [3] " exp ran successfully"
#> [4] " mean ran successfully"
```

9.4 The Maybe Monad: Handling Errors

Another common monad is **Maybe**, which handles computations that might fail. Instead of crashing, functions return either:

- `Just(value)` if successful
- `Nothing` if something went wrong

9 Robust Pipelines with Monads

The `{chronicler}` package uses this under the hood:

```
r_sqrt <- record(sqrt)

# This works
r_sqrt(16)
#>   Value: Just
#> [1] 4

# This fails gracefully
r_sqrt("not a number")
#>   Value: Nothing
```

When `Nothing` is passed to a decorated function, it immediately returns `Nothing`—the error propagates through the pipeline without crashing.

9.5 Monads in Python: talvez

The same concepts exist in Python. The `talvez` package provides a `Maybe` monad:

```
from talvez import maybe, just

@maybe()
def parse_int(s: str) -> int:
    return int(s)

@maybe(ensure=lambda x: x != 0)
def reciprocal(n: int) -> float:
    return 1 / n
```

```
# Successful computation
result = (
    parse_int("25")
        .bind(reciprocal)
        .fmap(lambda x: x * 100)
)
print(result)          # Just(4.0)
print(result.get_or(-1)) # 4.0

# Failed computation
bad = (
    parse_int("not a number")
        .bind(reciprocal)
        .fmap(lambda x: x * 100)
)
print(bad)  # Nothing
```

The key operations are:

- `fmap(fn)`: Apply a pure function to the value inside the monad
- `bind(fn)`: Apply a function that itself returns a monad

9.6 Building Robust Pipelines

The real power of monads becomes apparent when you combine them with pipeline orchestration. Consider a typical data pipeline:

9 Robust Pipelines with Monads

```
# Standard pipeline - fragile
raw_data |>
  basic_cleaning() |>
  recodings() |>
  filter_arrivals() |>
  make_monthly() |>
  make_plot()
```

If `recodings()` fails halfway through, the entire pipeline crashes. You get an error message, but no information about what succeeded before the failure.

Now imagine wrapping each function with `record()`:

```
library(chronicler)

# Robust pipeline with logging
r_basic_cleaning <- record(basic_cleaning)
r_recodings <- record(recodings)
r_filter_arrivals <- record(filter_arrivals)
r_make_monthly <- record(make_monthly)
r_make_plot <- record(make_plot)

result <- raw_data |>
  r_basic_cleaning() |>
  bind_record(r_recodings) |>
  bind_record(r_filter_arrivals) |>
  bind_record(r_make_monthly) |>
  bind_record(r_make_plot)

# Now you get:
# - The result (or Nothing if any step failed)
# - A complete log of which steps ran
```

```
# - Exactly where and why it failed  
read_log(result)
```

This pattern transforms a fragile script into a robust, observable pipeline.

9.7 Integrating with rixpress

You can combine the power of `{chronicler}` with `{rixpress}` for even more robust pipelines. The key insight is that your `user_functions` can use `record()` internally:

```
# functions.R  
library(chronicler)  
  
# Create recorded versions of your functions  
r_basic_cleaning <- record(function(data) {  
  data |>  
    select(contains("TIME"), contains("20")) |>  
    pivot_longer(cols = contains("20"),  
                 names_to = "date",  
                 values_to = "passengers")  
})  
  
r_recodings <- record(function(data) {  
  data |>  
    mutate(traj_meas = fct_recode(traj_meas, ...)) |>  
    mutate(passengers = as.numeric(passengers))  
})  
  
# Export a pipeline function that uses bind_record
```

9 Robust Pipelines with Monads

```
process_aviation_data <- function(raw_data) {  
  raw_data |>  
    r_basic_cleaning() |>  
    bind_record(r_recodings)  
}
```

Then in your `{rixpress}` pipeline:

```
library(rixpress)  
  
list(  
  rxp_r_file(  
    name = avia_raw,  
    path = "data/avia.tsv",  
    read_function = readr::read_tsv  
) ,  
  
  rxp_r(  
    name = processed_data,  
    expr = process_aviation_data(avia_raw),  
    user_functions = "functions.R"  
) ,  
  
  # ... more steps  
) |>  
  rxp_populate()
```

The `{rixpress}` derivation caches the result while `{chronicler}` provides logging and error handling within the step.

9.8 Python Integration with talvez

The same pattern works in Python pipelines:

```
# functions.py
from talvez import maybe, chain, just

@maybe()
def basic_cleaning(df):
    return df.dropna().reset_index(drop=True)

@maybe()
def recodings(df):
    df['category'] =
    ↵ df['category'].map(category_mapping)
    return df

@maybe()
def filter_arrivals(df):
    return df[df['tra_meas'] == 'Arrivals']

def process_data(raw_df):
    """Returns Just(result) or Nothing with error
       handling."""
    return chain(
        just(raw_df),
        basic_cleaning,
        recodings,
        filter_arrivals
    )
```

Then in your {rixpress} pipeline:

```
rxp_py(  
    name = processed_data,  
    expr = "process_data(raw_data)",  
    user_functions = "functions.py"  
)
```

9.9 Summary

Monads add a layer of robustness to your pipelines:

- **chronicler** (`record()`, `bind_record()`): Logging and error handling for R
- **talvez** (`@maybe()`, `bind()`, `chain()`): Error handling for Python
- **Just/Nothing**: Graceful failure propagation without crashes
- **Composability preserved**: Even with extra capabilities, functions still compose

Combined with `{rixpress}`:

- Nix provides **hermetic, cached execution**
- Monads provide **logging and error handling**
- Together they give you **robust, reproducible, observable pipelines**

This completes our toolkit for reproducible data science: environments (Nix), functional code (FP), pipelines (rixpress), distribution (Docker), and robustness (monads).

References

Peng, Roger D. 2011. “Reproducible Research in Computational Science.” *Science* 334 (6060): 1226–27.

